

DataStax

Developers

SKY TRAINING

Introduction to Cassandra

Building efficient applications with Apache Cassandra



DataStax Sky account team

James Kenney - Account Manager

Amarjit Basra - Data Architect

Matt Chalmers - PS manager

Astyn Gould - Astra Success Team

Sky Slack Channel : #datastax-sky-shared-channel (95 members)

Face-to-Face meetings at Osterley Site



Your DataStax team



Director of Developer Relations



- Trainer
- Public Speaker
- Developers Support
- Developer Applications
- Developer Tooling

- Creator of ff4j (ff4j.org)
- Maintainer for 8 years+

- Happy developer for 14 years
- Spring Petclinic Reactive & Starters
- Implementing APIs for 8 years



clunven



clunven



clun



Cédrick Lunven



Solution Architect



DataStax
Enterprise

PULSAR



Google Cloud

Azure

IBM Cloud



kubernetes



mongoDB



ORACLE
DATABASE

- Joined DataStax in January 2022
- 28 years in Enterprise Software
- Wide industry experience

- Support Developers and Enterprises in creating and optimising Cassandra & DSE Solutions

- Back-end technology "junkie"



AIX

hpUX

ORACLE
Solaris

PURESTORAGE



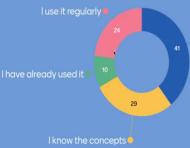
DataStax

Phil Miesle

sky

Games and quizzes: menti.com

How much experience do you have with the Spring Framework ?



Mentimeter®

Questions: <https://dtsx.io/discord>

Screenshot of a Discord channel titled "workshop-chat". The channel has over 100 members. A message from user "louis901" asks if it's ok to use a background image for the workshop. Another user, "Eric Farnier", suggests adding some programming background to help understand the concepts. A third user, "louis901", asks about transitioning to machine learning. A user named "louis901" also asks for help with Apache Cassandra.



Discord
(#workshop-chat)

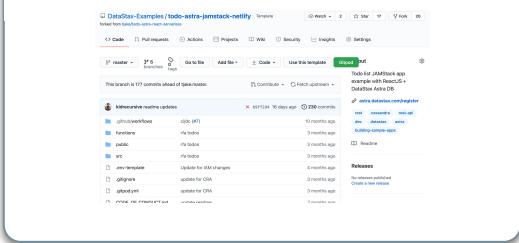


DataStax

Interactive Workshops

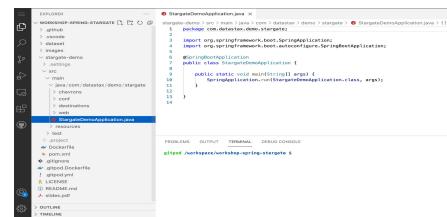
Nothing to install !

Source code + exercises + slides



!github

IDE



!gitpod

Database + Api + Streaming



DataStax
Astra DB

!astra

Hands-On Housekeeping

572 and counting



Intro to Cassandra Workshop

Awarded to **YOUR NAME** • attendee@workshop.com

Issued on Jun 16, 2021

Offered By
[DataStax Developers](#)

Upgrade Complete! This badge is to certify successful completion of the DataStax Cassandra Workshop: "Fundamentals of Apache Cassandra".



Verified
Last verified by Badgr on Jun 17, 2021

[Re-verify Badge](#)

EARNING CRITERIA

Recipients must complete the earning criteria to earn this Badge

To earn this badge, individuals must complete the following during the [Intro to Cassandra Workshop](#):

- Attend the lecture
- Complete the practical steps by doing all required



Achievement Unlocked!

01



Why NoSQL ?

02

Introduction to
Apache Cassandra™

03

Distributed Architecture
The CAP Theorem

04

Tables
and Partitions

05

The Art of
Data Modelling

06

What's next?
Quiz, Homework, Next week

DataStax

DataStax Developers

Agenda

sky

01



Why NoSQL ?

02

Introduction to
Apache Cassandra™

03

Distributed Architecture
The CAP Theorem

04

Tables
and Partitions

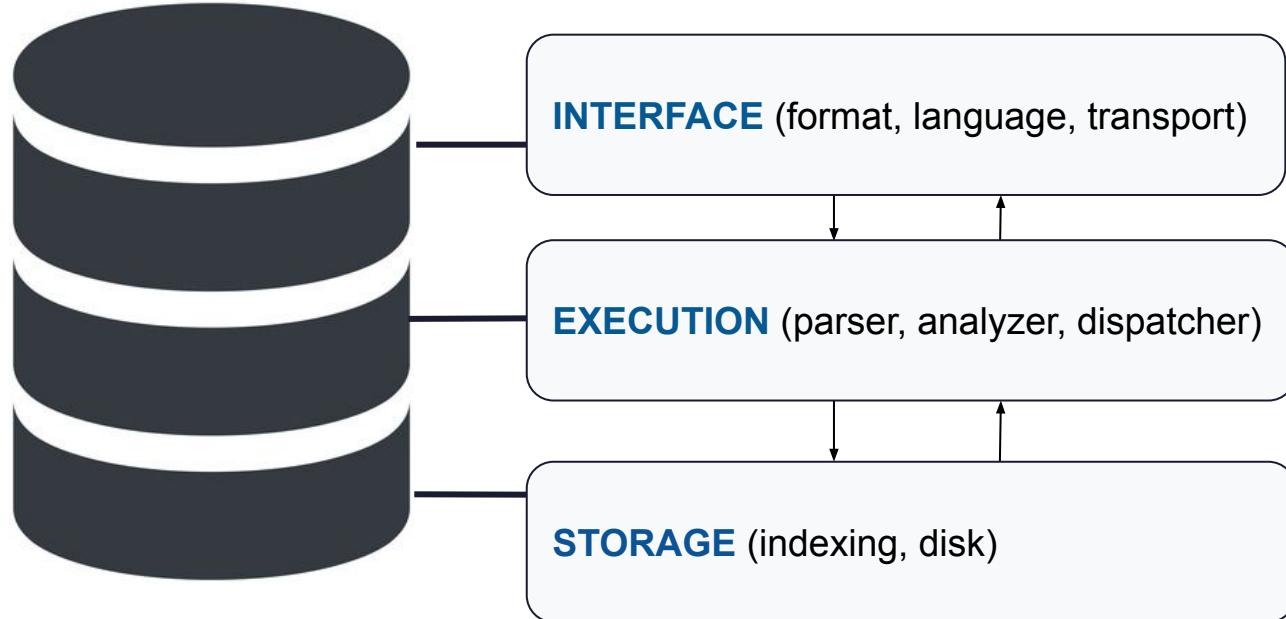
05

The Art of
Data Modelling

06

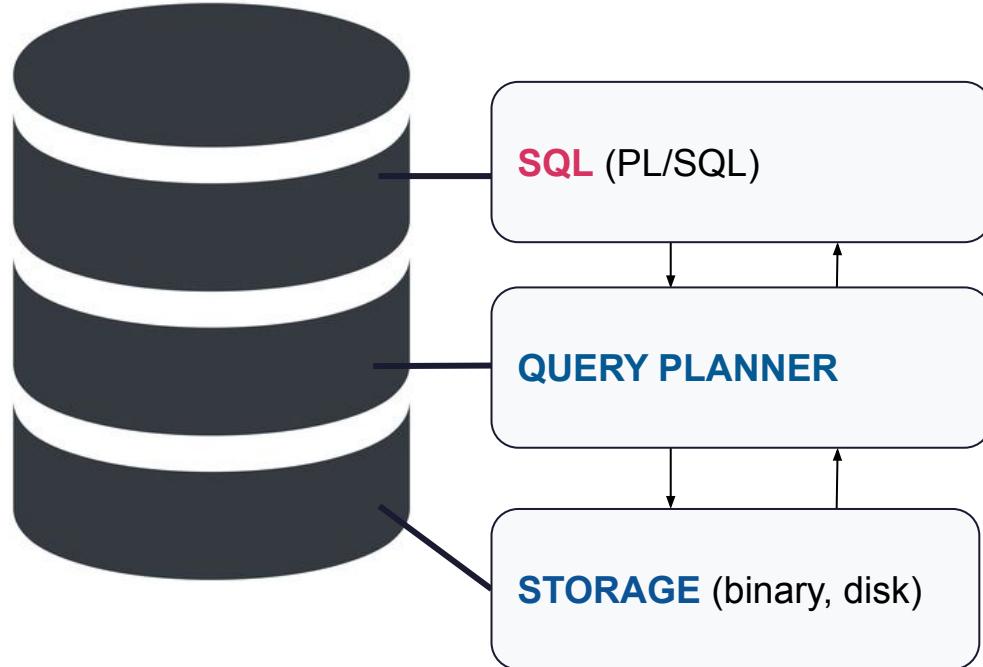
What's next?
Quiz, Homework, Next week

DataStax



Introducing NoSQL





Introducing NoSQL

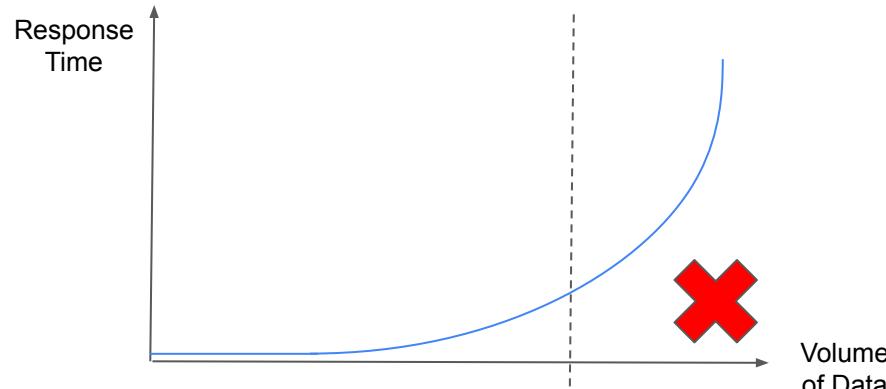


OLTP: Online Transaction Processing

Fast Processing
Transactional
High number of transactions
Hot / Live Data
“Row oriented”
Normalized Data (3NF)

OLAP: Online Analytical Processing

Slow Queries
Historical
High volume of data
Cold Data
“Column oriented”
Denormalized Data



Relational Databases are **versatile**



- They have been designed to run on a single machine

- New requirements

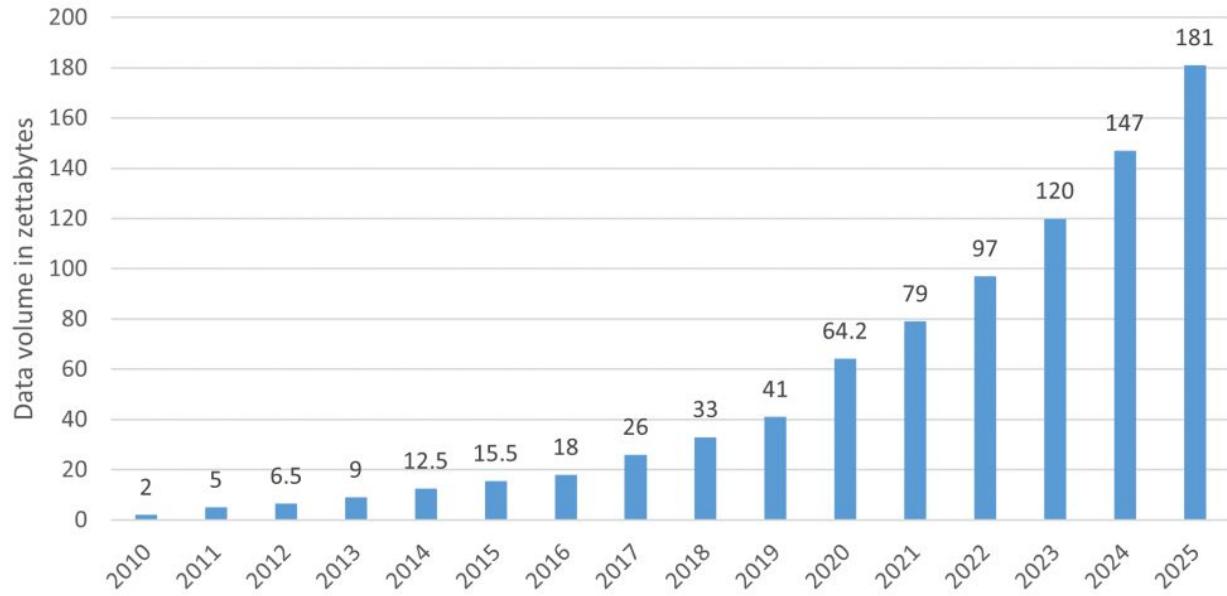
- V: VOLUME
- V: VELOCITY
- V: VARIETY

WHAT'S A ZETTABYTE?	
1 kilobyte	1,000 000,000,000,000,000,000
1 megabyte	1,000,000 000,000,000,000,000
1 gigabyte	1,000,000,000 000,000,000,000
1 terabyte	1,000,000,000,000 000,000,000
1 petabyte	1,000,000,000,000,000 000,000
1 exabyte	1,000,000,000,000,000,000,000
1 zettabyte	1,000,000,000,000,000,000,000,000

SOURCE: CISCO



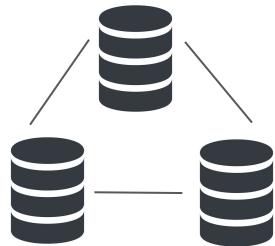
Volume of data created and replicated worldwide (source: IDC)



Relation Databases have **limited scalability**



- Meetup names on June 11, 2009 in San Francisco
- Web Giants needed more scalable systems
 - Distributed by design (horizontal scalability)
 - Data is replicated
 - Cope with the variety of data



NoSQL



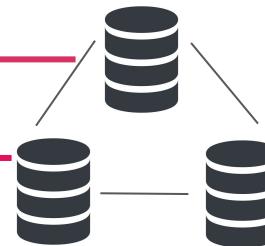
DataStax

Introducing NoSQL



AP vs CP (CA = not distributed)

Server vs VM vs Container vs Cloud



SQL or JSON or CQL or N1QL or Cypher...

Query Parser and execution engine

Text, Binary, Graph

Ledger Databases

awsqldb, ksql_db

Time-Series Databases

Influx, OpenTSDB, Prometheus

Tabular Databases

Cassandra, Hbase, Bigtable

Object Databases

S3, Minio, Ceph

Document Databases

MongoDB, Elastic, DocDb

Graph Databases

Neo4j, Datastax Graph, Titan

Key/Value Databases

Redis, Dynamo, Memcache

DataStax

DataStax Developers

NoSQL Databases

sky

01



Why NoSQL ?

02

Introduction to
Apache Cassandra™

03

Distributed Architecture
The CAP Theorem

04

Tables
and Partitions

05

The Art of
Data Modelling

06

What's next?
Quiz, Homework, Next week



Agenda



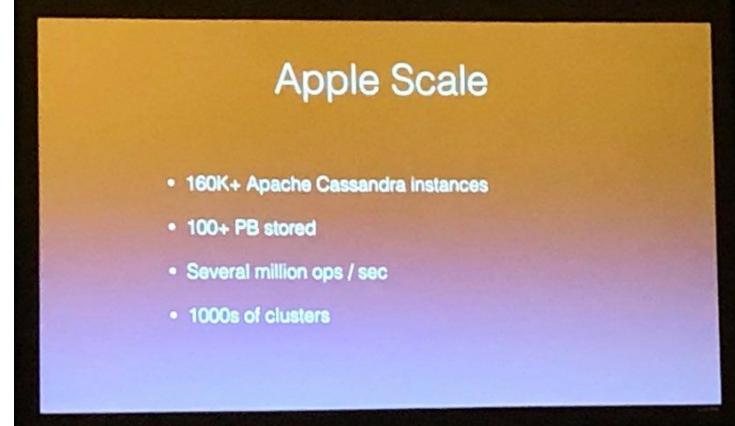
Apache Cassandra @ Netflix

- . 98% of streaming data is stored in Apache Cassandra
- . Data ranges from customer details to viewing history to billing and payments
- . Foundational datastore for serving millions of operations per second

- 30 million ops/sec on most active single cluster
- 500 TB most dense single cluster
- 9216 CPUs in biggest cluster

O(100) Clusters
O(10000) Instances
O(10,000,000) Replications per second
O(100,000,000) Operations per second
O(1,000,000,000,000,000) Petabytes of data

dtsx.io/cassandra-at-netflix



And many others...



Cassandra Biggest Users (and Developers)

Apache Cassandra

NoSQL Distributed Decentralised Database Management System



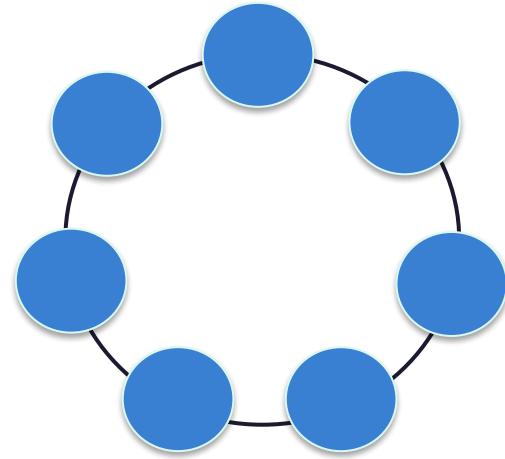
- Big Data Ready
- Read / Write Performance
- Linear Scalability
- Highest Availability
- Self-Healing and Automation
- Geographical Distribution
- Platform Agnostic
- Vendor Independent



Cassandra Features



Partitioning over distributed architecture makes the database capable to handle data of any size: we mean petabytes scale. Need more volume? Add more nodes.

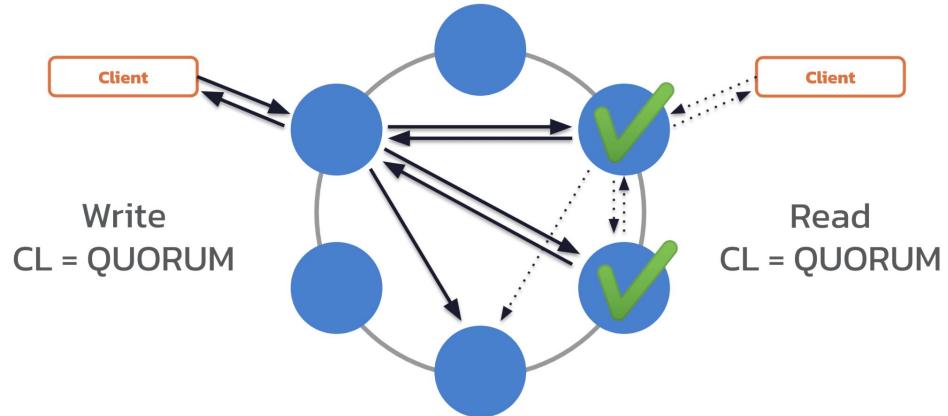


Big Data Ready



Even a single Cassandra node is very performant but a cluster consisting of multiple nodes and data centers brings throughput to the next level.

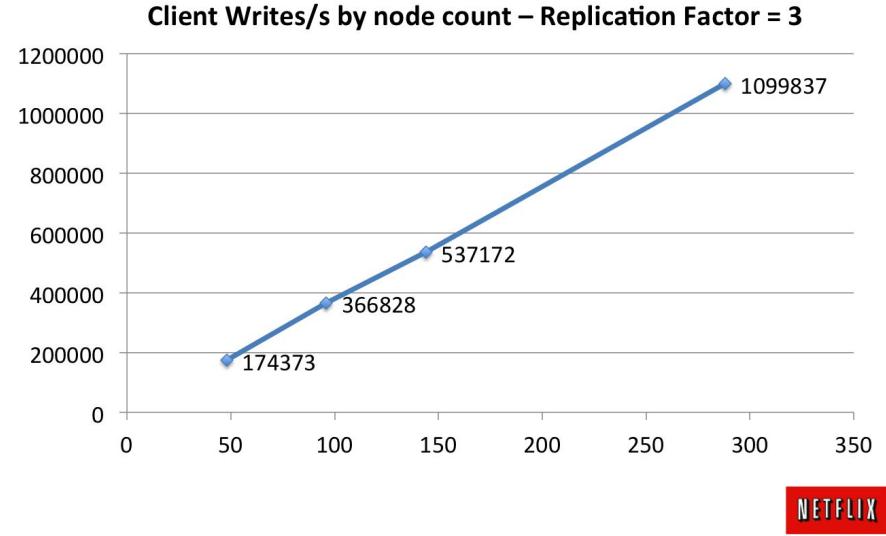
Decentralisation (**masterless architecture**) means that every node is able to deal with any request, read or write.



Read / Write Performance



- For volume or velocity, there are no limitations
- **Linear** - No overhead on new nodes, scales with your needs*

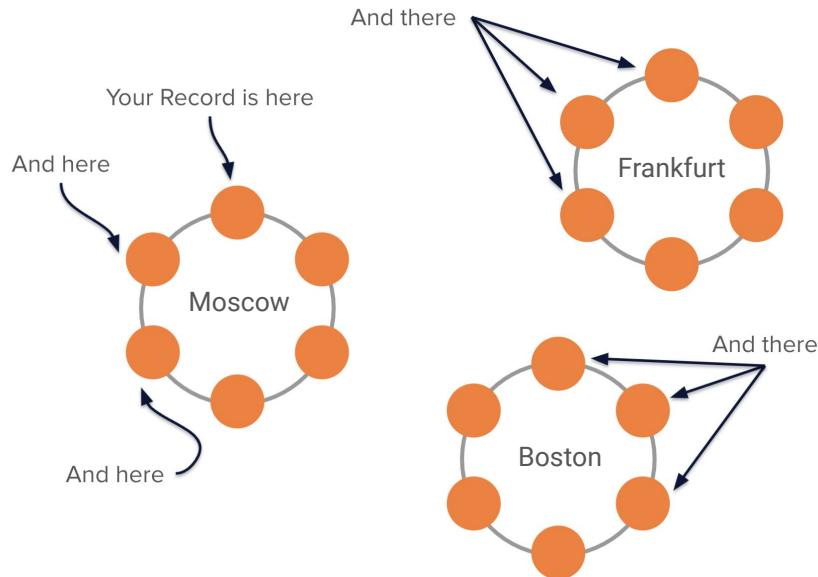


Linear Scalability



Replication, Decentralisation, and Topology-Aware Placement Strategy take care of possible downtimes:

- Multiple Live Replicas
- No Single Point of Failure
- Network topology-aware data placement
- Client-side Smart Reconnection and Strong Retry Mechanism



Highest Availability



Operations for a huge cluster can be exhausting so Apache Cassandra clusters are smart and able to scale, change data placement and recover automatically.

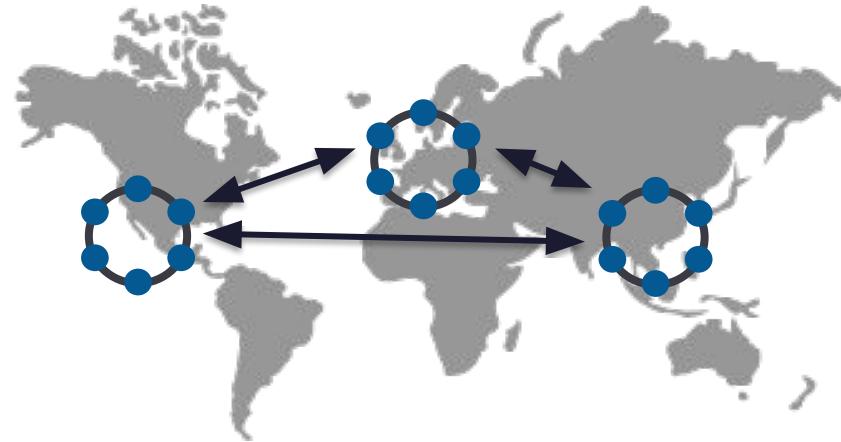


Self-Healing and Automation



Cassandra's trademark is multi-datacenter deployments, granting you an exceptional capability for disaster tolerance while keeping your data close to your clients - worldwide.

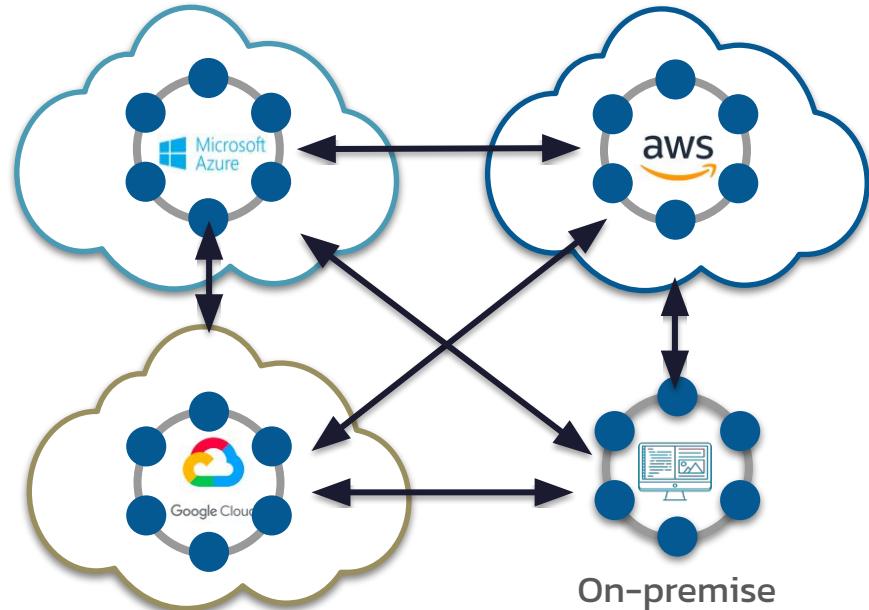
All DCs are active (available for both writes and reads)!



Geographical Distribution



Apache Cassandra is **not bound to any platform** or service provider, helping you build hybrid-cloud and multi-cloud solutions with ease.



Platform Agnostic

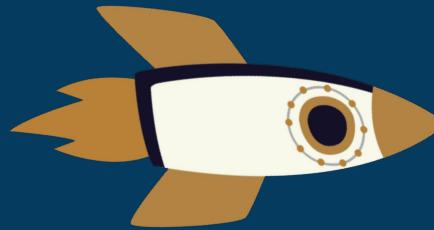


Cassandra doesn't belong to any of commercial vendors but controlled by a non-profit Open Source **Apache Software Foundation**, already familiar to you by *Hadoop*, *Spark*, *Kafka*, *Zookeeper*, *Maven* and many other projects.



Vendor Independent





Let's GO (!github)

-  1. Objectives
-  2. Frequently asked questions
-  3. Materials (Slides)





Free to Use

Up to 80GB storage and/or 20 million operations monthly.



Serverless

Lower your costs by running Cassandra clusters only when needed.



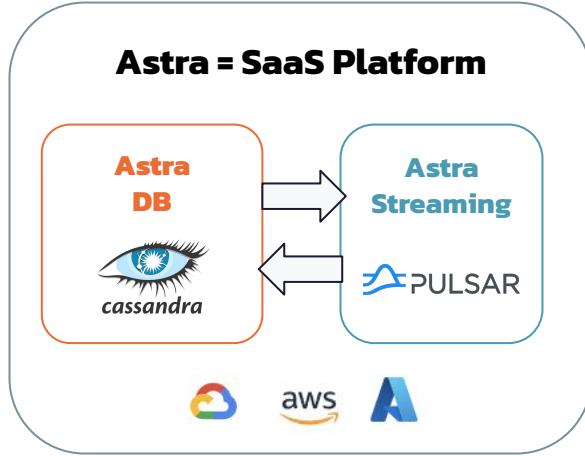
No Operations

Eliminate the overhead to install, operate, and scale Cassandra.



Data APIs

Work natively with Document (JSON), REST, GraphQL and gRPC APIs.



Global Scale

Put your data where you need it without compromising performance, availability or accessibility.



End-to-End Security

Secure connect with VPC peering and Private Link. Bring your own encryption key management. SAML SSO secure account access.



Zero Lock-in

Deploy on AWS, GCP or Azure and keep compatibility with open-source Cassandra.



Relational Indexes

Storage Attached Indexing (SAI) lets you query tables using any columns.



Astra = Cassandra As a Service++





Lab 1

Database Setup

[http://github.com/datastaxdevs/
workshop-cassandra-fundamentals](http://github.com/datastaxdevs/workshop-cassandra-fundamentals)

- ✓ Create a new database
- ✓ Wake up an existing hibernated database

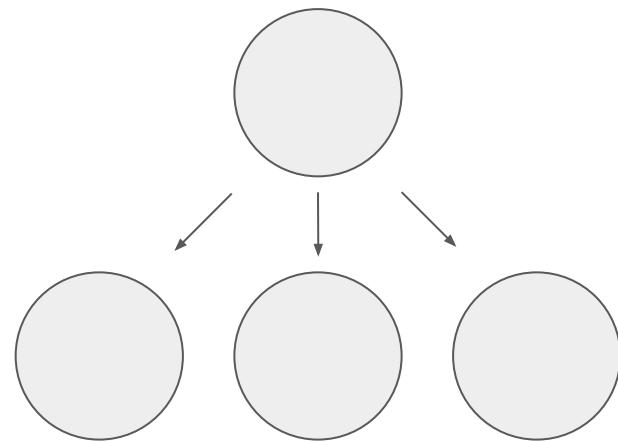


All servers are created equal

Cassandra Holy Book, P.I Paragraph I



Leader Server (Write/Read)

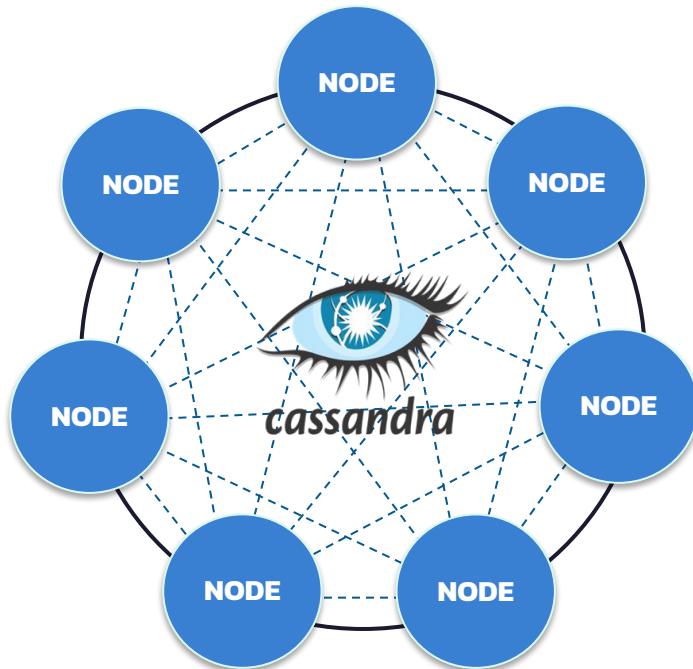


1. Single Point of Failure
2. Hard to scale for writes
3. Application needs to know where to write



Traditional Architecture





1. NO Single Point of Failure
2. Scales for writes and reads
3. Application can contact any node
(in case of failure - just contact next one)



Master-less (Peer-to-Peer) Architecture

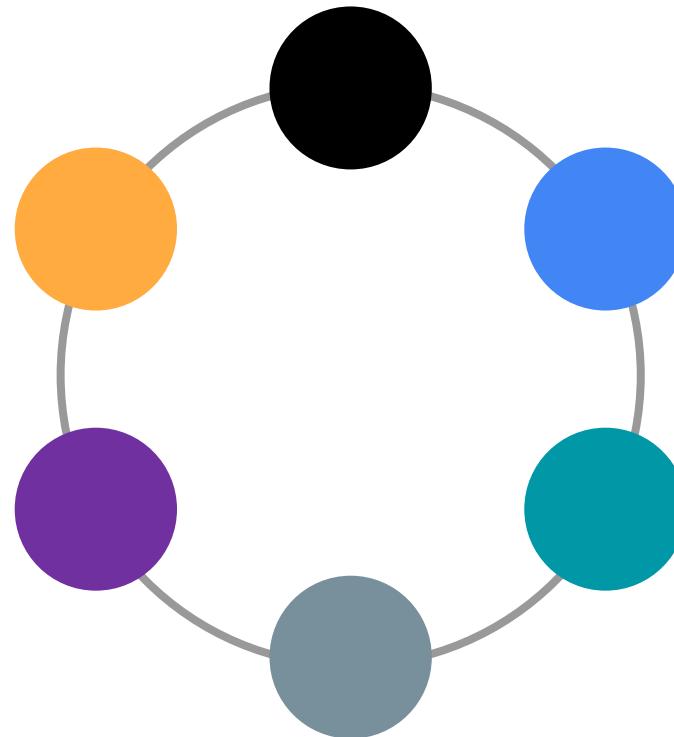


Data is Replicated



RF = ?

Replication Factor
means the number
of nodes used to
store each partition



Replication Factor



```
CREATE KEYSPACE sensor_data  
  WITH REPLICATION = {  
    'class' : 'NetworkTopologyStrategy',  
    'us-west-1' : 3,  
    'eu-east-2' : 5  
};
```

keyspace replication strategy



Replication factor by data center



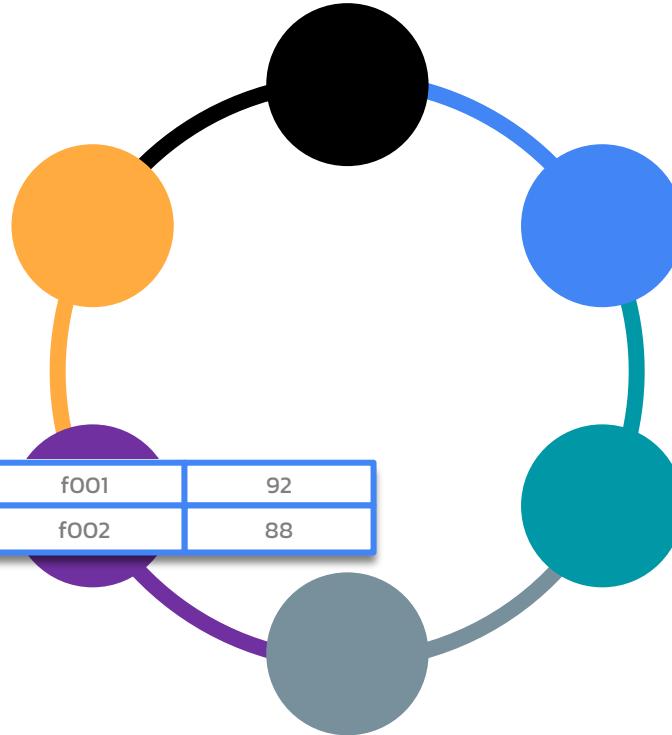
Data is Replicated



RF = 1

Replication Factor 1
means that every
partition is stored
on 1 node

forest	f001	92
forest	f002	88

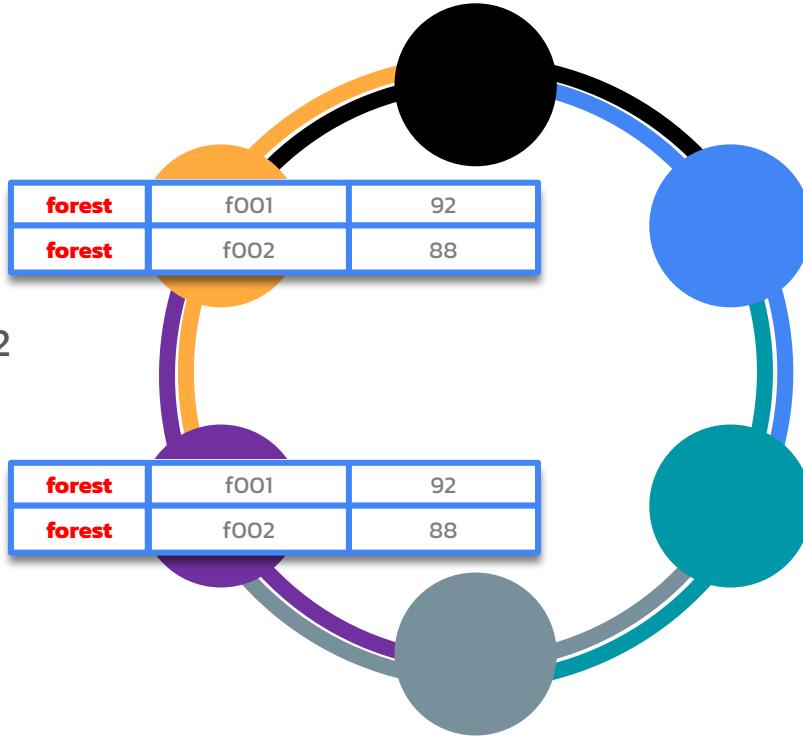


Replication Factor



RF = 2

Replication Factor 2
means that every
partition is stored
on 2 nodes

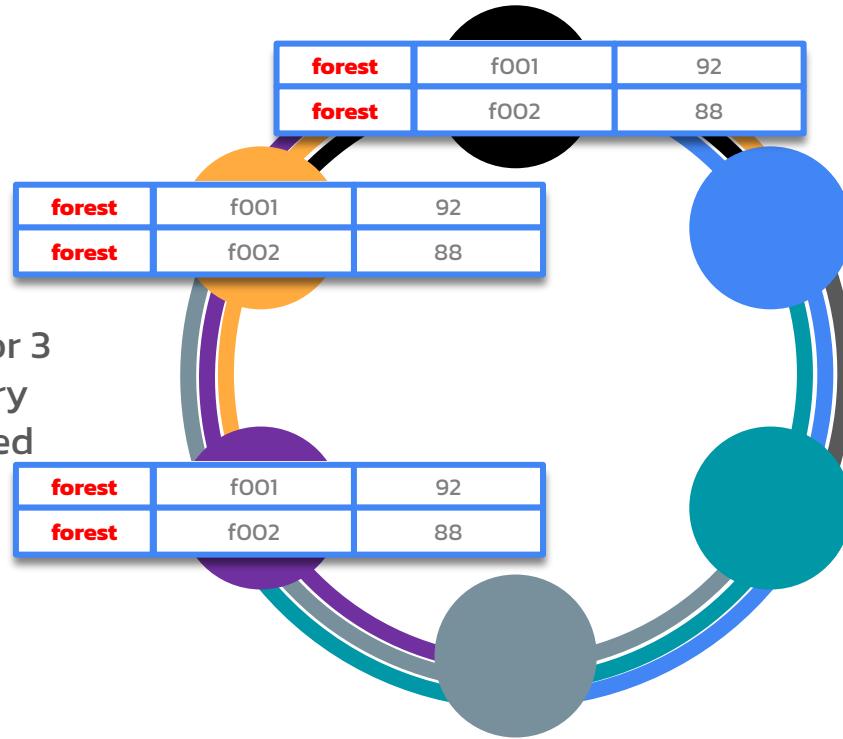


Replication Factor



RF = 3

Replication Factor 3
means that every
partition is stored
on 3 nodes

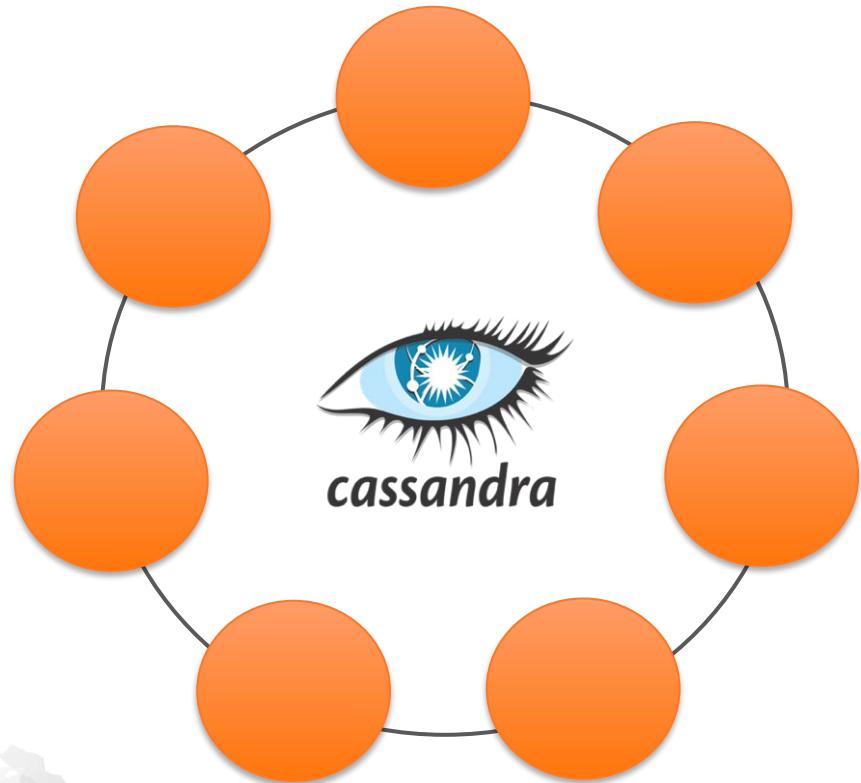


Replication Factor



Data is Distributed





Data is Distributed

network	sensor	temperature
forest	f001	92
forest	f002	88
volcano	v001	210
sea	s001	45
sea	s002	50
home	h001	72
car	c001	69
car	c002	70
dog	d001	40
road	r001	105
road	r002	110
ice	i001	35

Partition Key



forest	f001	92
forest	f002	88

Network

Sensor

temperature

sea	s001	45
sea	s002	50

car	c001	69
car	c002	70

volcano	v001	210
----------------	------	-----

home	h001	72
-------------	------	----

ice	i001	35
dog	d001	40

road	r001	105
road	r002	110



cassandra



Data is Distributed



```
CREATE TABLE sensor_data.sensors_by_network (
    network      text,
    sensor       text,
    temperature integer,
    PRIMARY KEY ((network), sensor)
);
```

Keyspace

Table

Primary key

Partition key

Clustering columns



Key-based Partitioning



network	sensor	Value
forest	f001	92
forest	f002	88
sea	s001	45
sea	s002	50
road	r001	105
road	r002	110

Partition Keys

Partitioner
Murmur3 Hashing

Network	Sensor	Value
59	f001	92
59	f002	88
12	s001	45
12	s002	50
45	r001	105
45	r002	110

Tokens

Cassandra Nodes



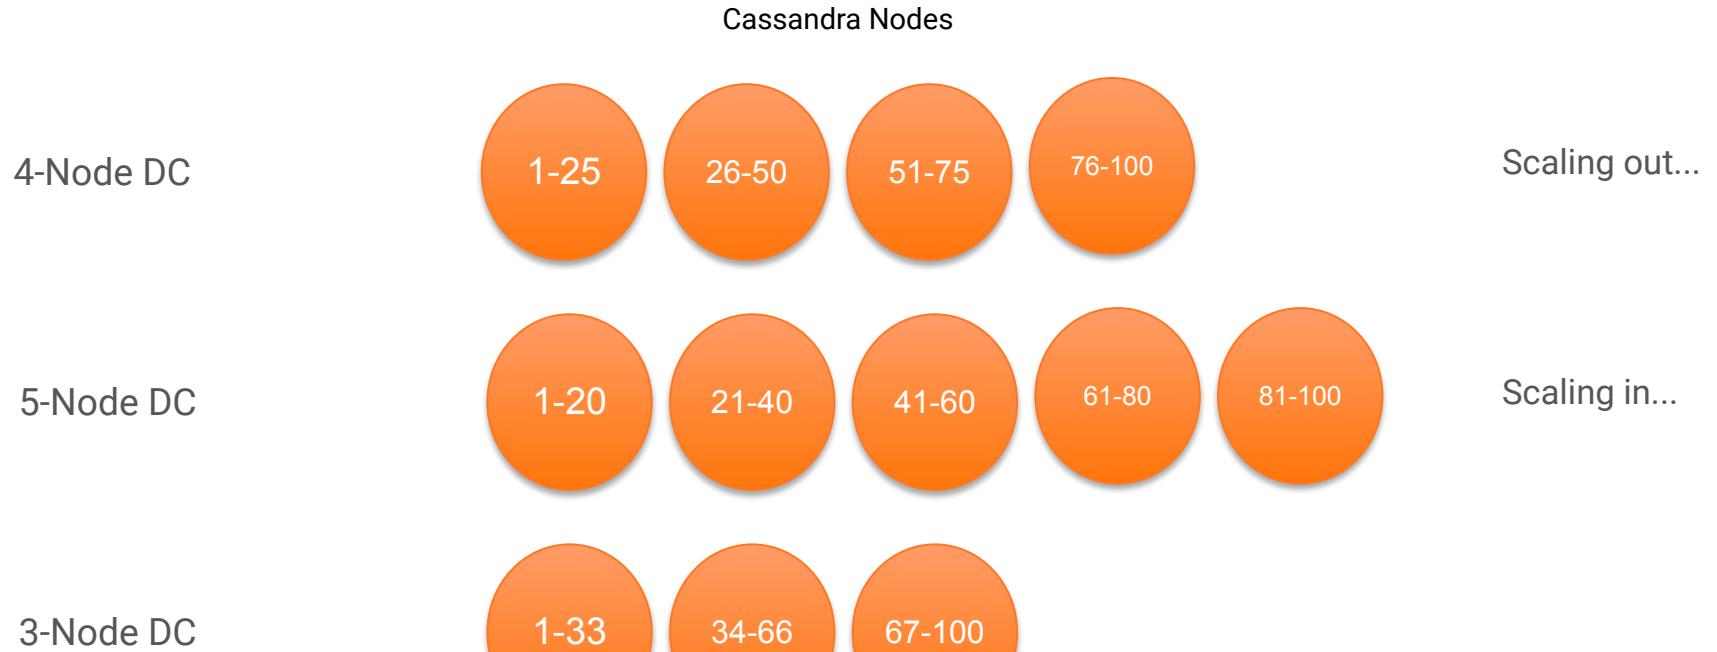
Partitioning and Token Ranges



Why partitioning?
Because scaling doesn't have to be [s]hard!

Big Data doesn't fit to a single server, splitting it into chunks we can easily spread them over dozens, hundreds or even thousands of servers, adding more if needed.





NOTICE: No downtime! All scaling operations are live!



Token Ranges Recalculation

Replication Factor = 1

1-25

26-50

51-75

76-100

Replication Factor = 2

1-25,
26-50

26-50,
51-75

51-75,
76-100

76-100,
1-25

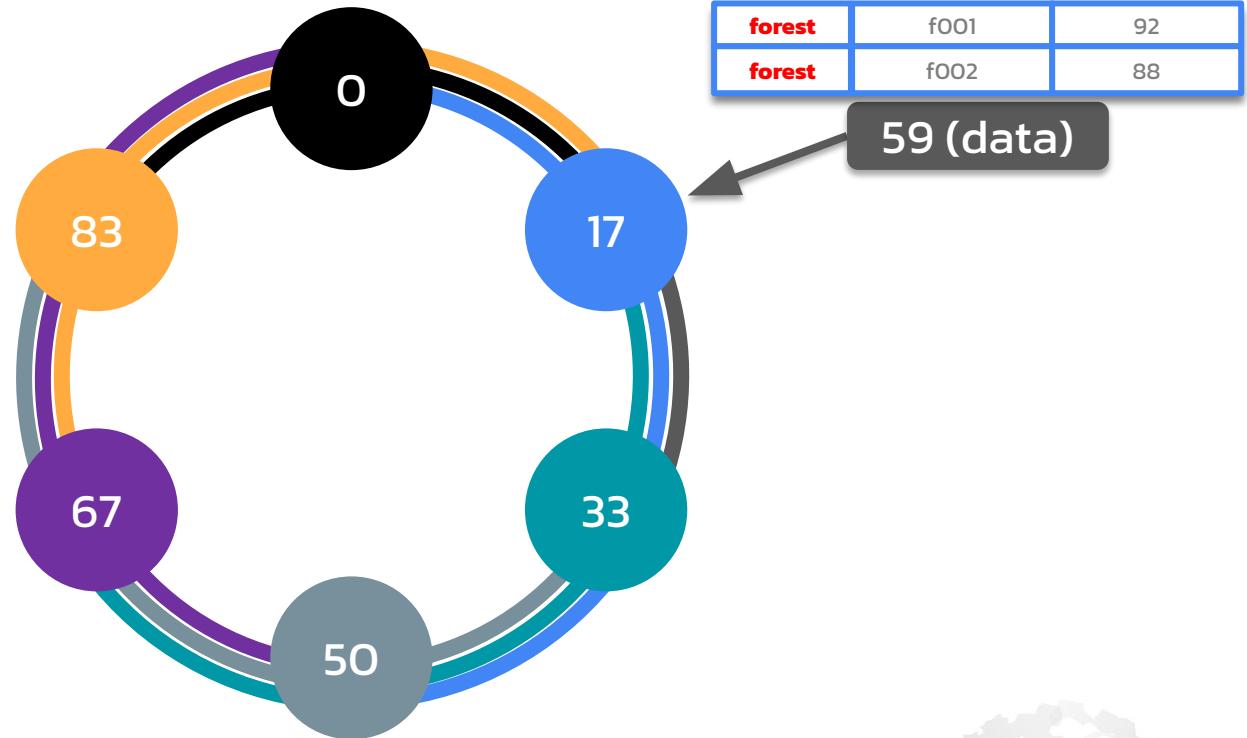
And so on...



Partitioning + Replication



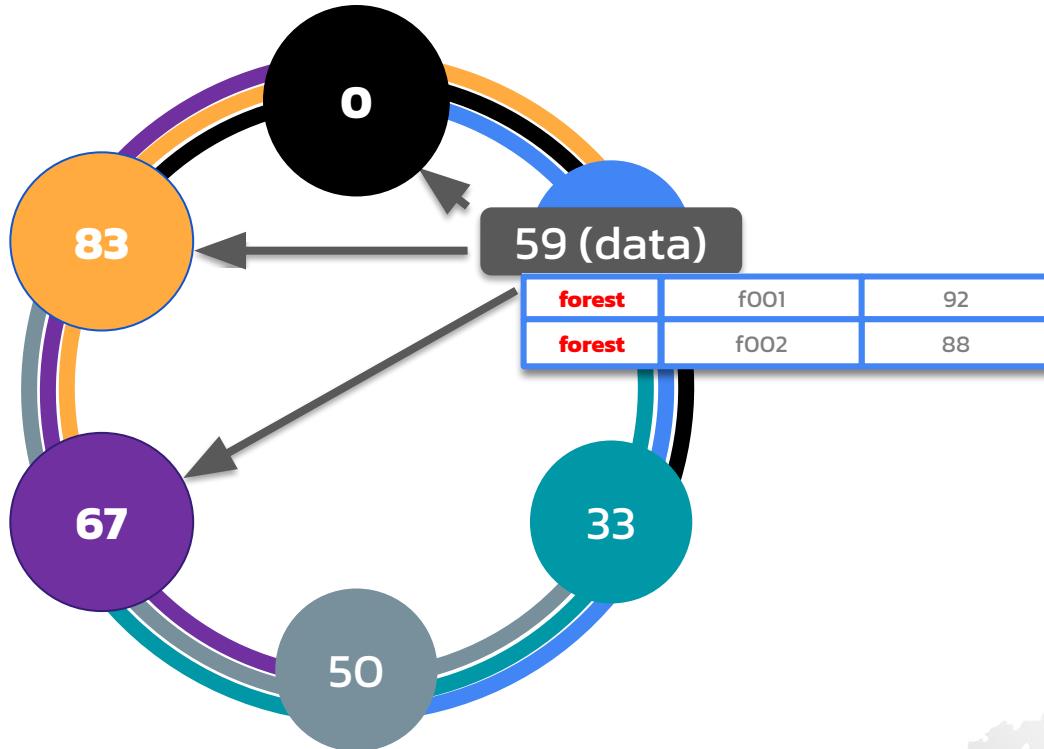
RF = 3



Partitioning + Replication



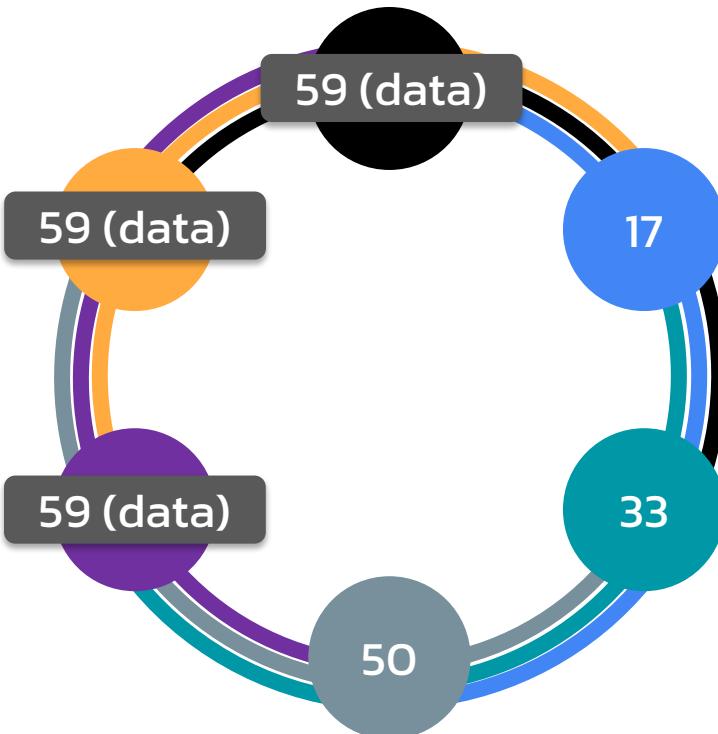
RF = 3



Partitioning + Replication



RF = 3



Partitioning + Replication



IMPORTANT

Each Cassandra node and even each Cassandra driver knows Data Allocation in a cluster (it's called Token-Aware), so your application can contact literally ANY server and still get the answer.

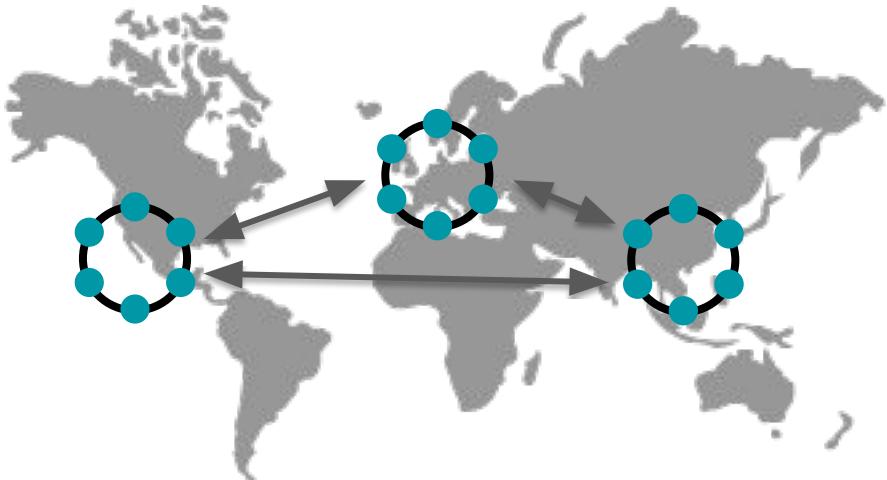




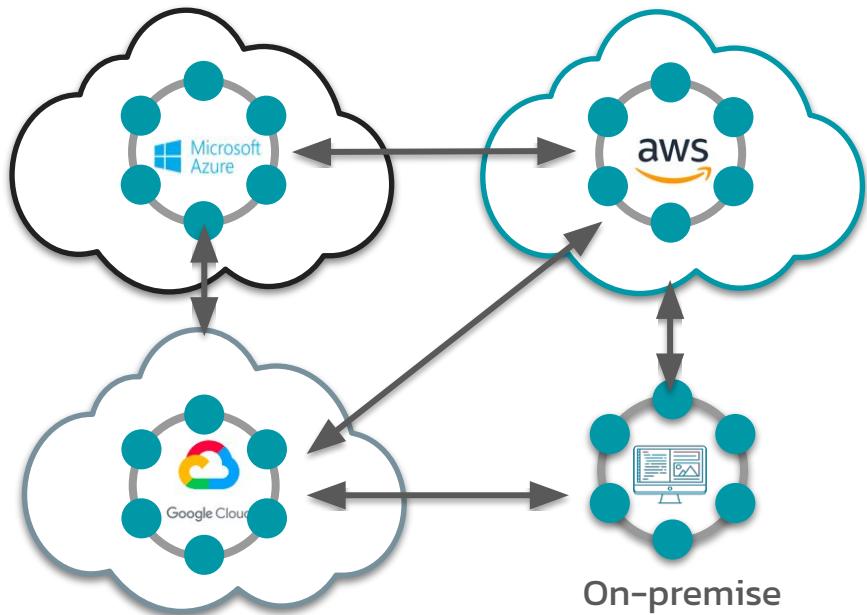
Data is GLOBALLY distributed



Geographical Distribution



Hybrid-Cloud and Multi-Cloud

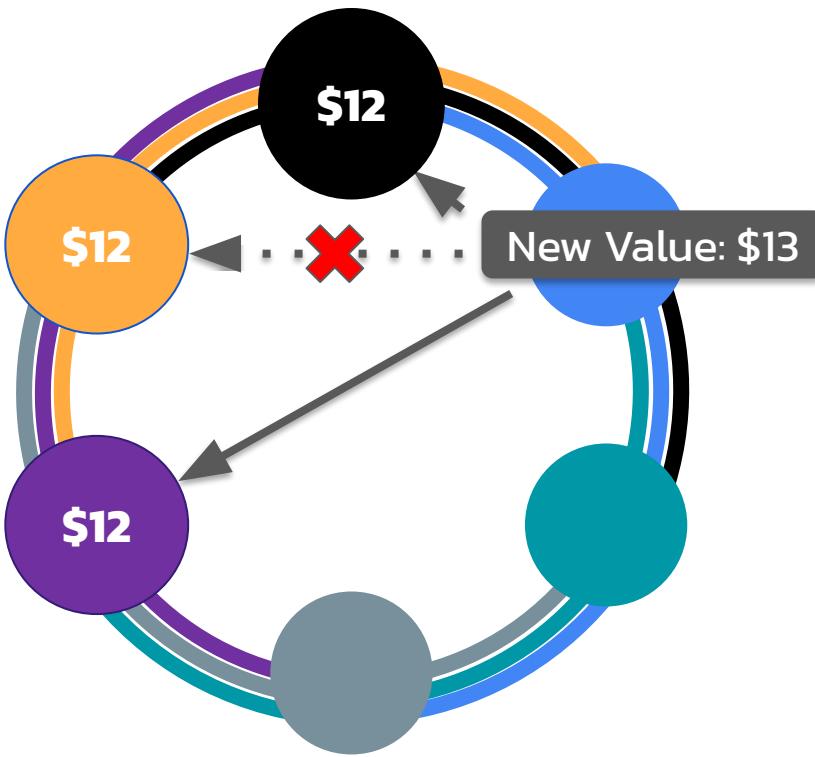


Data is globally distributed



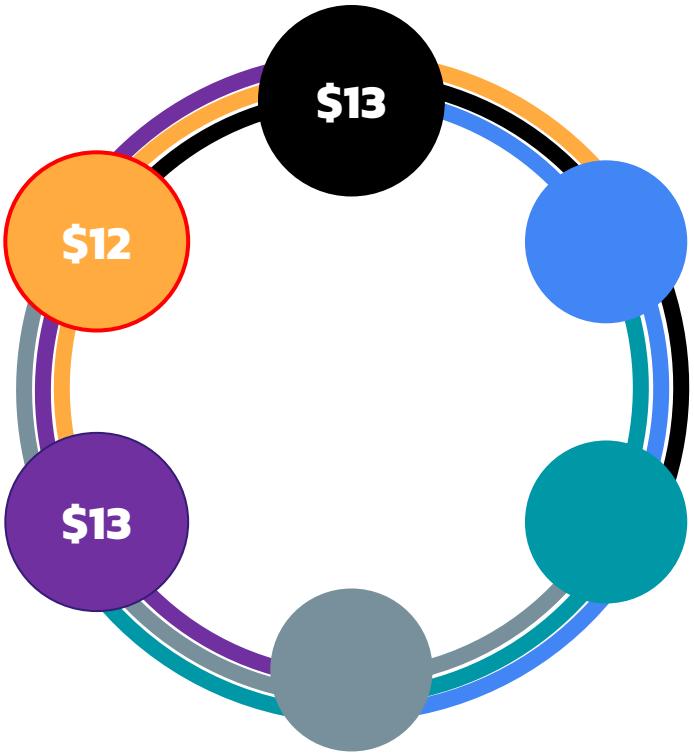


What is the biggest problem of replication?



Potential Inconsistency



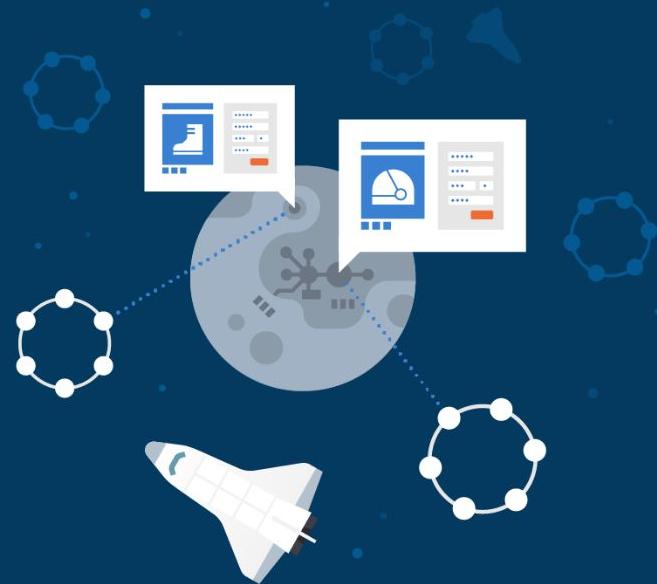


DataStax Developers

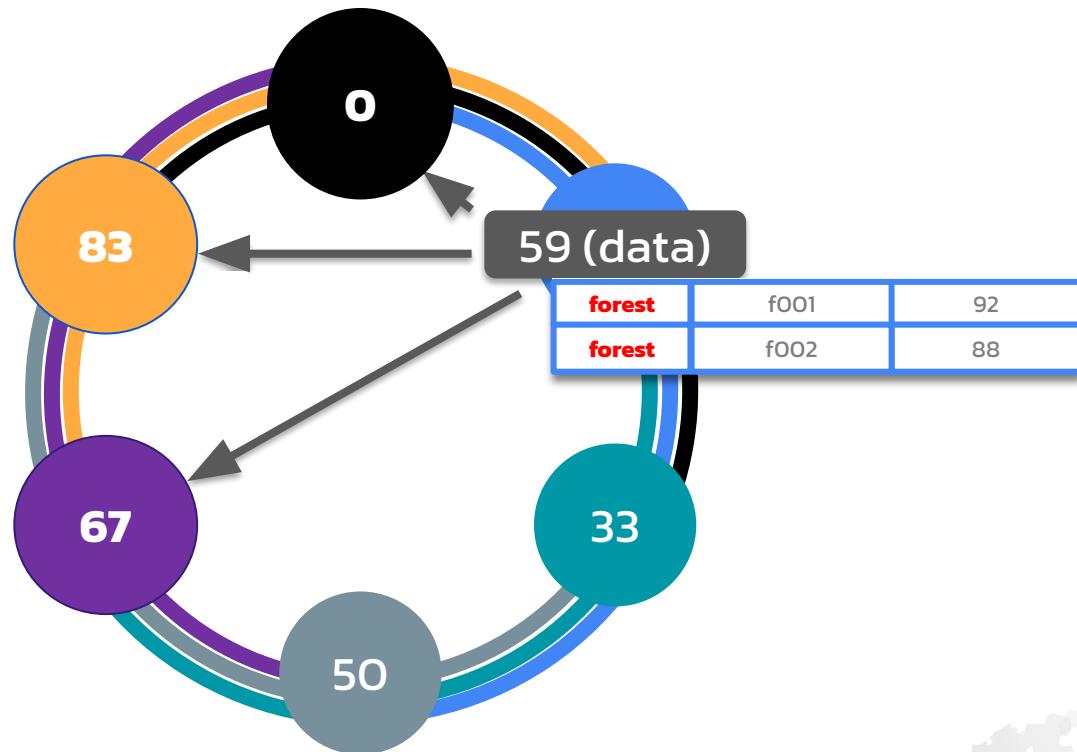
Potential Inconsistency



Cassandra Layered Self-Defence



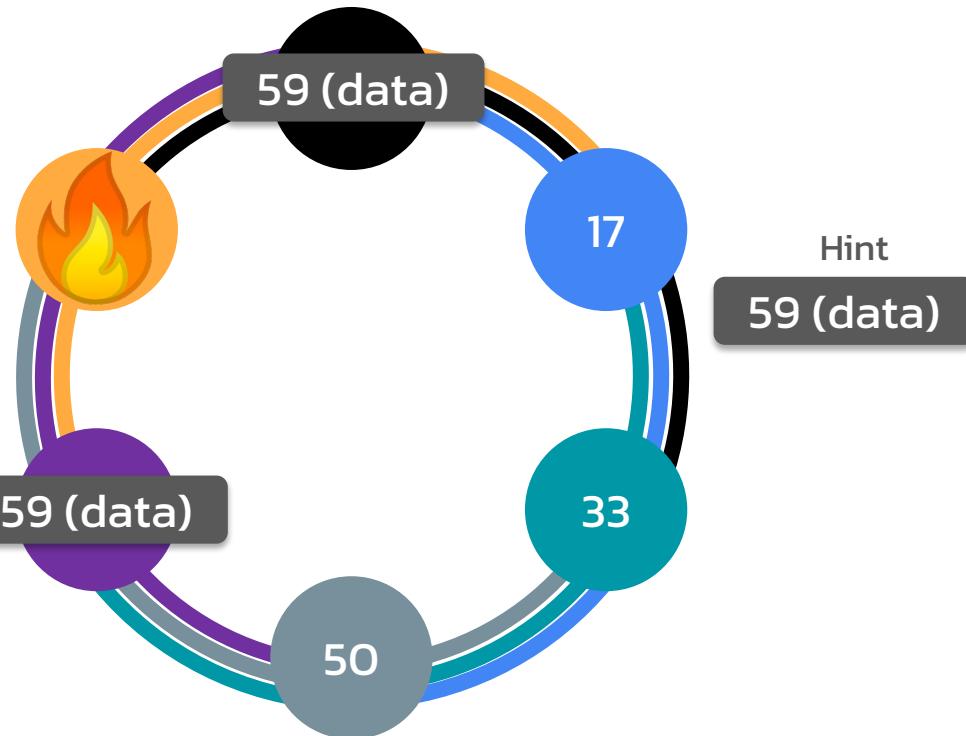
RF = 3



Hinted Hand-offs



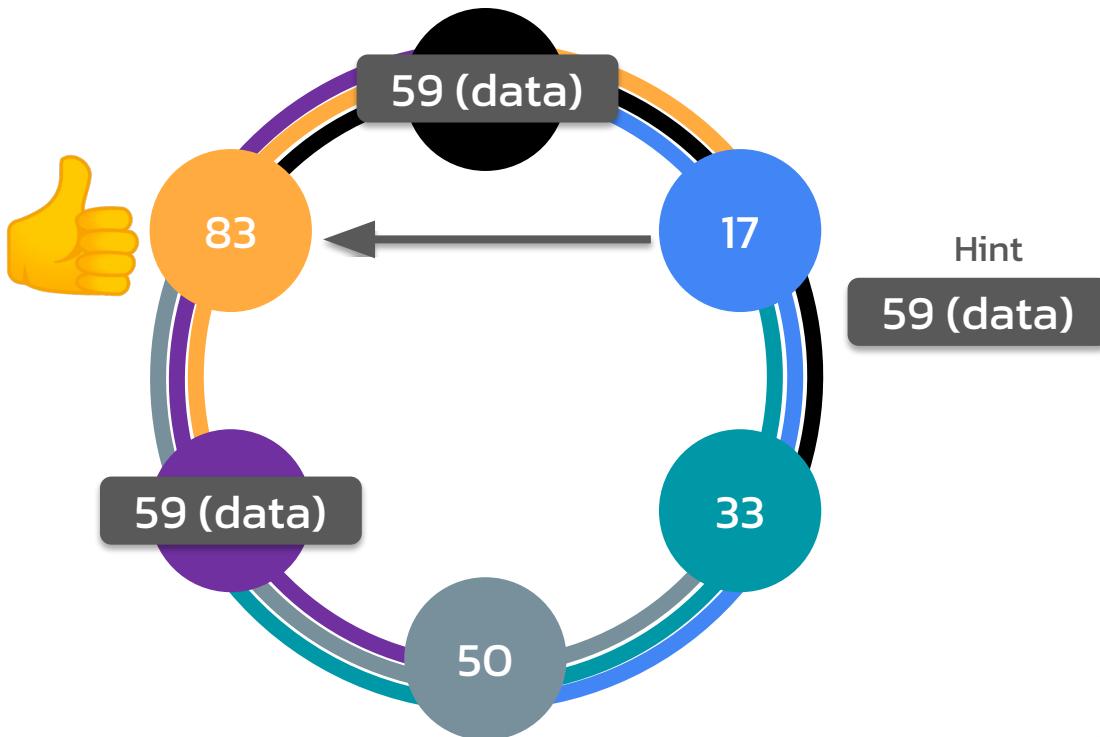
RF = 3



Hinted Hand-offs



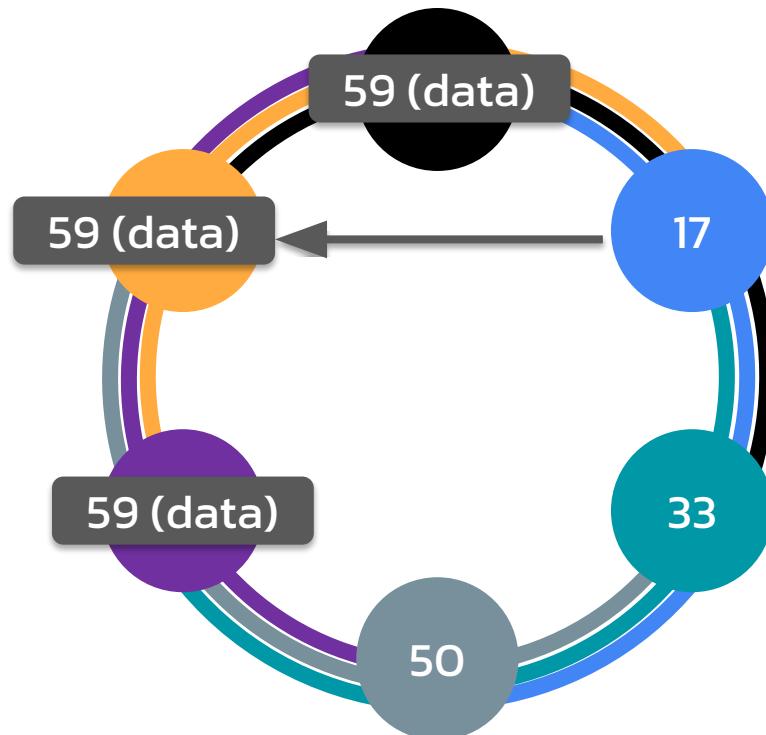
RF = 3



Hinted Hand-offs



RF = 3



Hinted Hand-offs



01



Why NoSQL ?

02

Introduction to
Apache Cassandra™

03

Distributed Architecture
The CAP Theorem

04

Tables
and Partitions

05

The Art of
Data Modelling

06

What's next?
Quiz, Homework, Next week

DataStax

DataStax Developers

Agenda

sky

CAP Theorem operates three features:

1. Availability
2. Consistency
3. Partition Tolerance



CAP Theorem Fundamentals



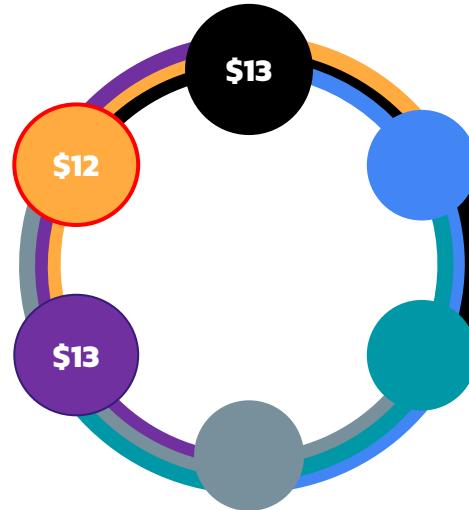
Availability basically means “Uptime”. You ask the question, you get the answer. If failure of a single or even multiple servers doesn’t lead to no response (no downtime), your system is **available**.



CAP Theorem: Availability



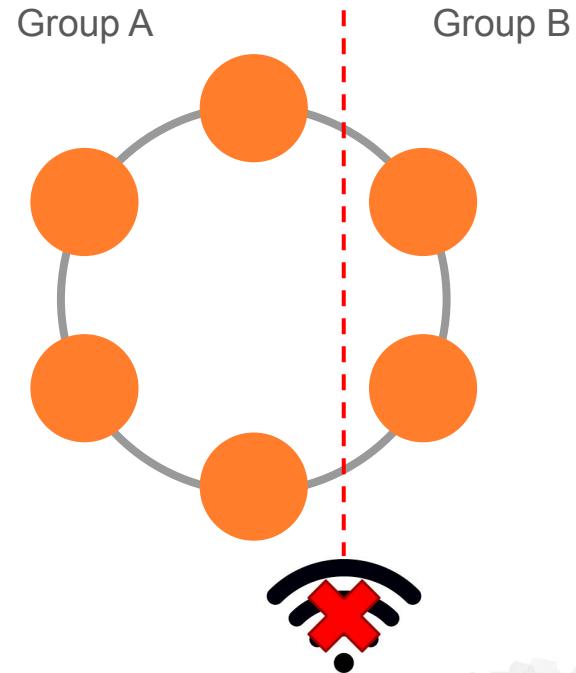
Consistency means “no stale data”. You ask for something, you get the most recent value. If one of your servers return outdated information, your system is **inconsistent**.



CAP Theorem: Consistency



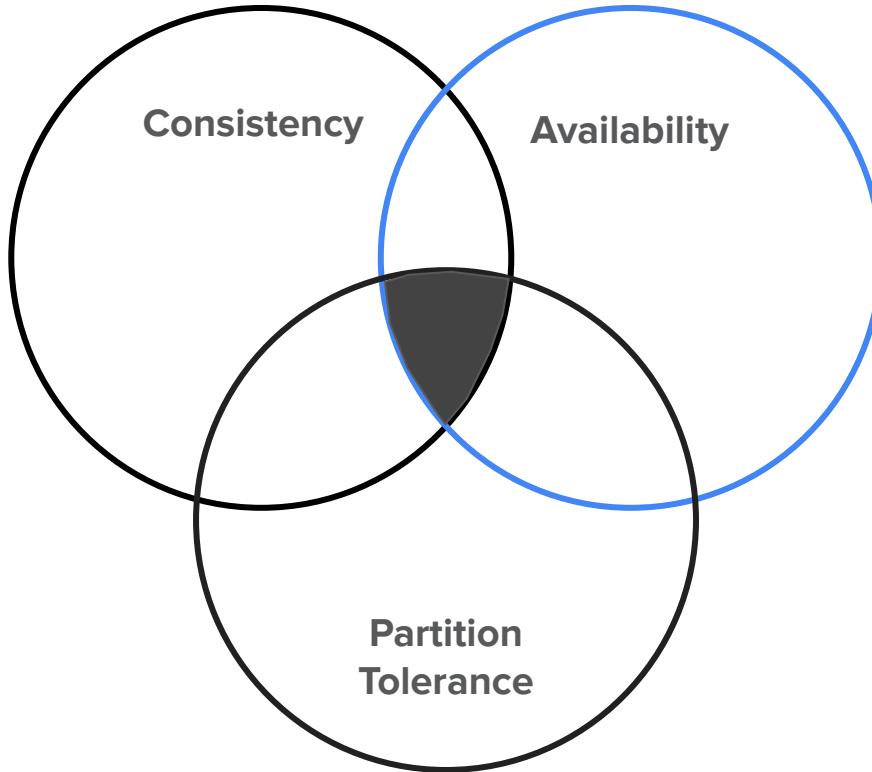
Partition Tolerance is the ability of a distributed system to survive “network partitioning”. Network partitioning means that the part of the servers can not reach the second part.



CAP Theorem: Partition Tolerance



In the distributed environment **in case of emergency** you can have only two guaranteed qualities out of three :(



CAP Theorem



Cassandra is configurable consistent. In any moment of the time, for any particular query you can set the Consistency Level you require to have. It defines how many **CONFIRMATIONS** you'll wait before the response is dispatched;

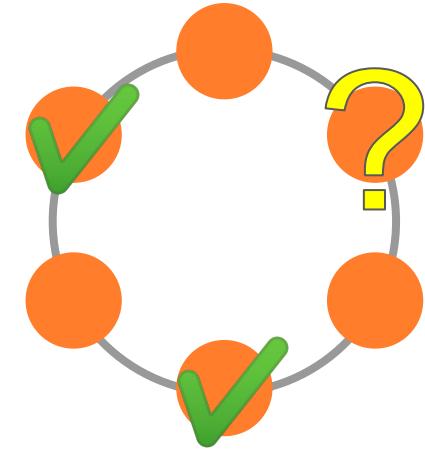
```
PreparedStatement pstmt = session.prepare(  
    "INSERT INTO product (sku, description) VALUES (?, ?)"  
);  
pstmt.setConsistencyLevel(ConsistencyLevel.ONE);
```

```
cqlsh> CONSISTENCY
```

```
Current consistency level is QUORUM.
```

```
cqlsh> CONSISTENCY ALL
```

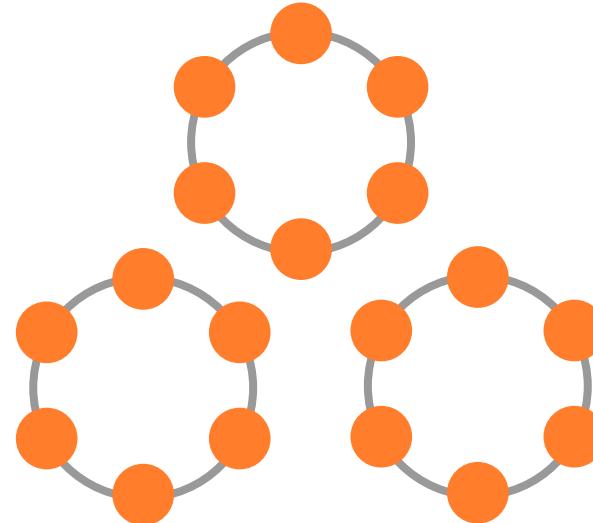
```
Consistency level set to ALL.
```



Is Cassandra AP or CP?

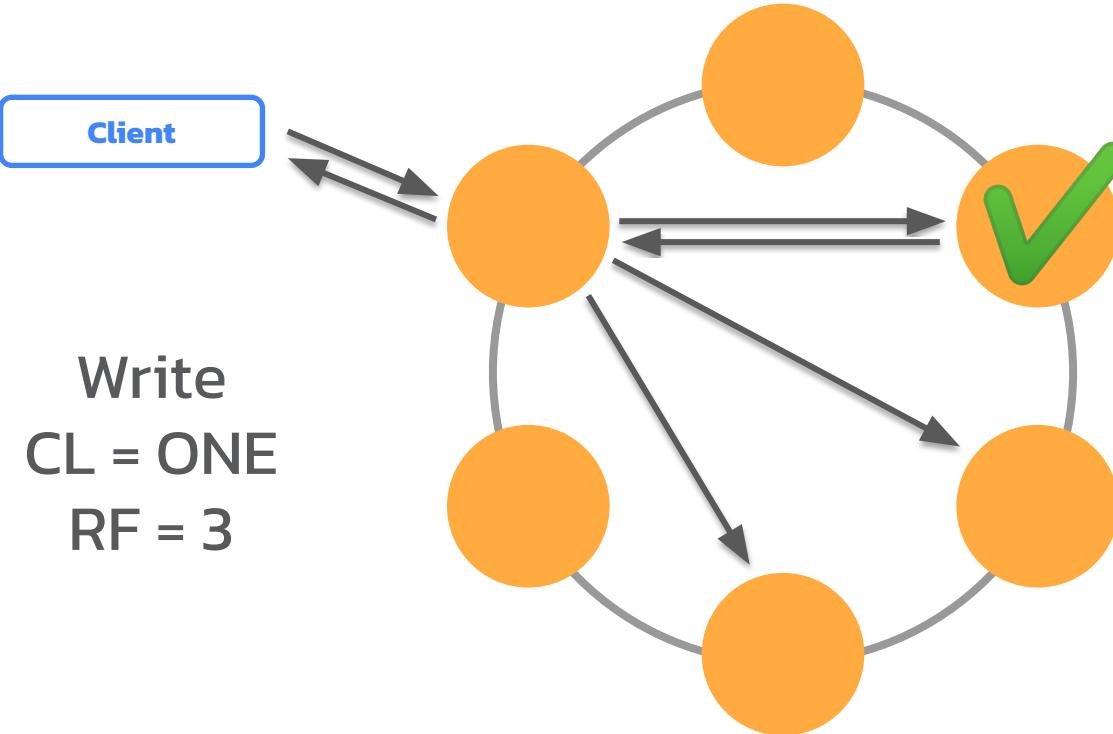


- ANY
- ONE
- LOCAL_ONE
- TWO, THREE
- QUORUM
- LOCAL_QUORUM
- EACH_QUORUM
- ALL



Query Consistency Levels



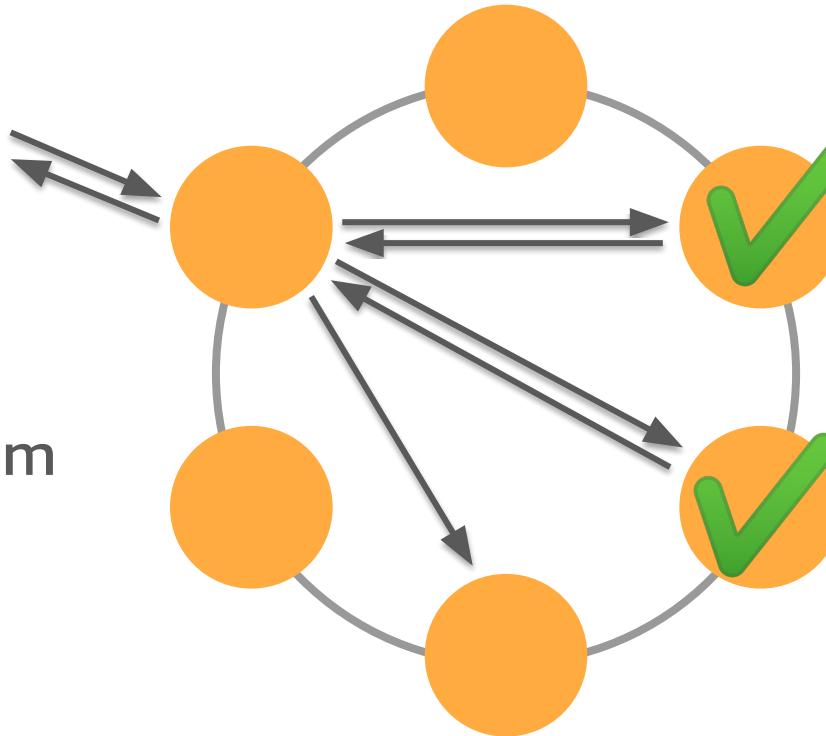


Consistency Level One



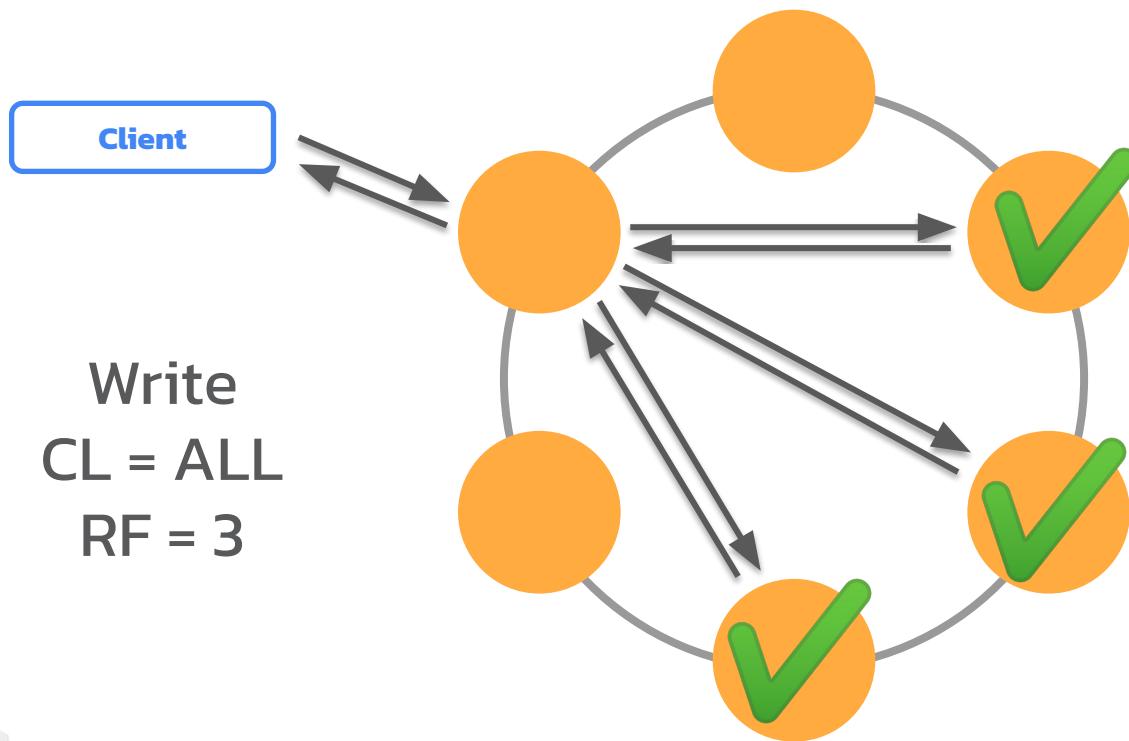


Write
CL = Quorum
RF = 3



Consistency Level Quorum





Write
CL = ALL
RF = 3



Consistency Level ALL



Client

Writes
CL = ALL
RF = 3

IT'S A TRAP!

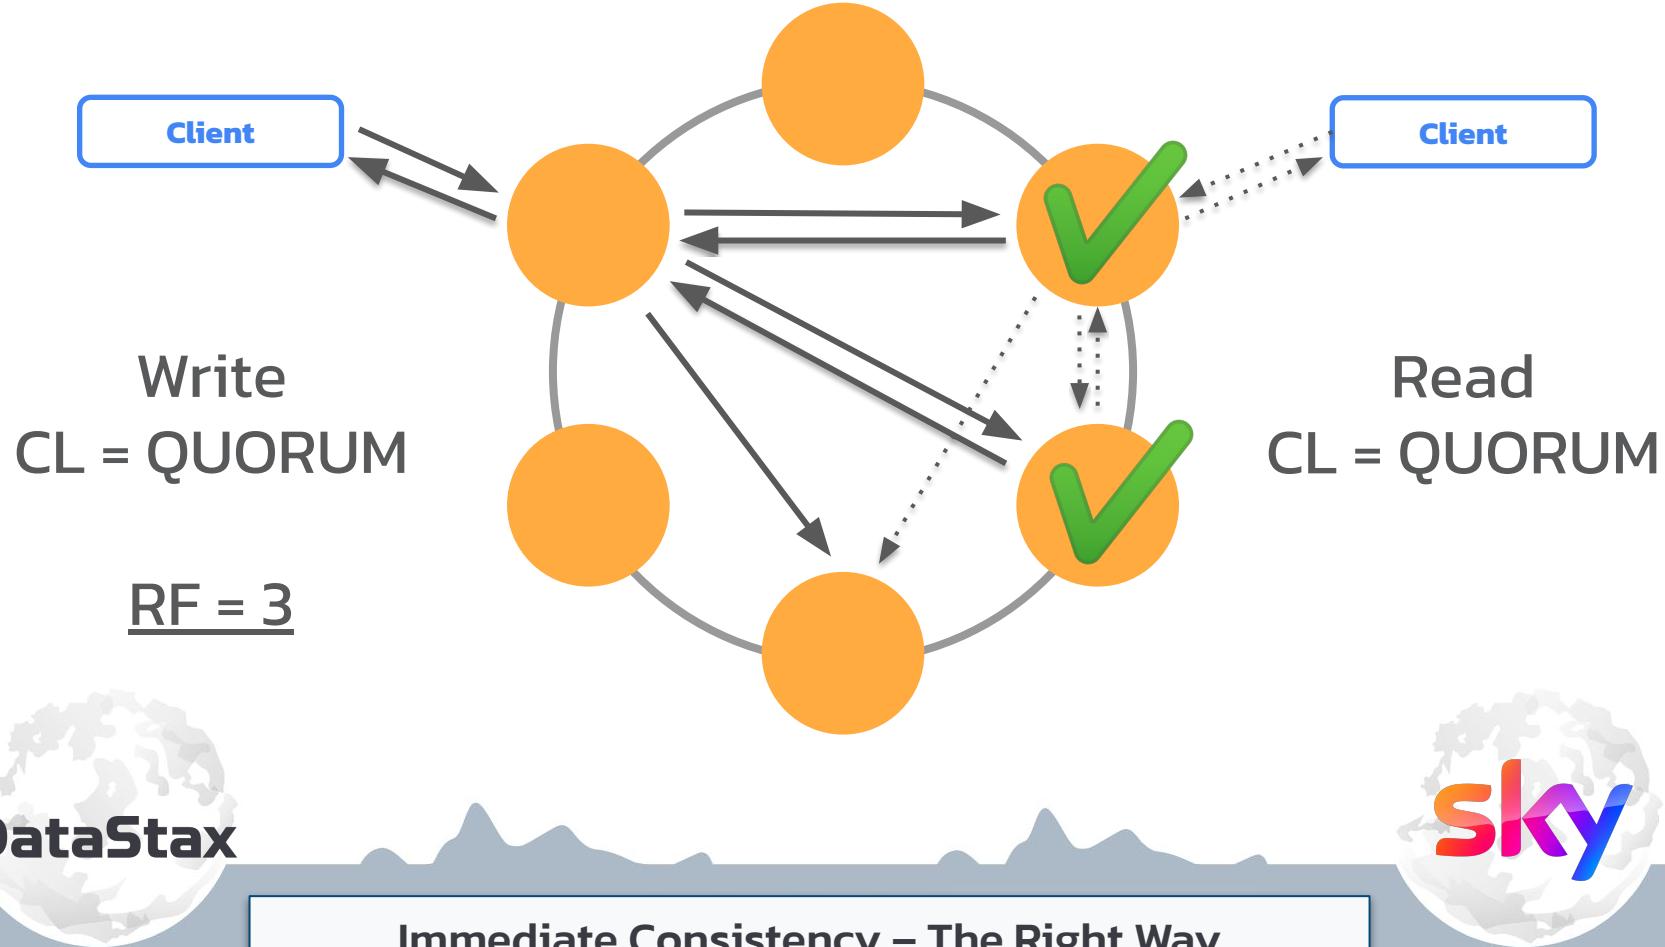
CAP THEOREM IS STILL HERE!



Consistency Level ALL



CL Write + CL Read > RF → Immediate Consistency



01



Why NoSQL ?

02

Introduction to
Apache Cassandra™

03

Distributed Architecture
The CAP Theorem

04

Tables
and Partitions

05

The Art of
Data Modelling

06

What's next?
Quiz, Homework, Next week

DataStax

Data Structure: a Cell



An intersection of a row
and a column, stores data.

foo1



Data Structure: a Cell



Data Structure: a Row



A single, structured data item in a table.

forest	f001	92
---------------	-------------	-----------



Data Structure: a Row



Data Structure: a Partition



A group of rows having the same partition token, a base unit of access in Cassandra.

IMPORTANT: stored together, all the rows are guaranteed to be neighbors.

forest	f001	92
forest	f002	88
forest	f003	90
sea	s001	45



Data Structure: a Partition



Data Structure: a Table



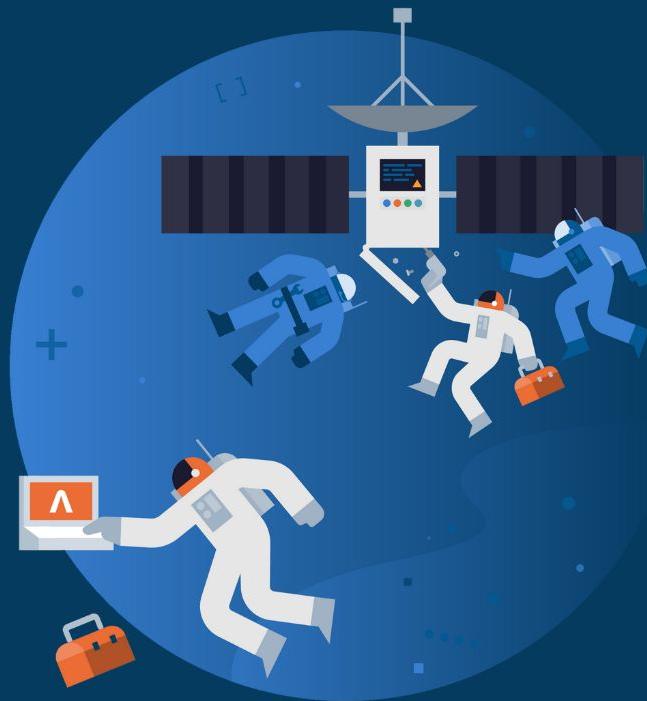
A group of columns and rows storing partitions.

network	sensor	temperature
forest	f001	92
forest	f002	88
volcano	v001	210
sea	s001	45
sea	s002	50
home	h001	72
car	c001	69
car	c002	70
dog	d001	40
road	r001	105
road	r002	110
ice	i001	35



Data Structure: a Table





Lab 2

Create Tables

[http://github.com/datastaxdevs/
workshop-cassandra-fundamentals](http://github.com/datastaxdevs/workshop-cassandra-fundamentals)



01



Why NoSQL ?

02

Introduction to
Apache Cassandra™

03

03

Distributed Architecture
The CAP Theorem

04

Tables
and Partitions

05

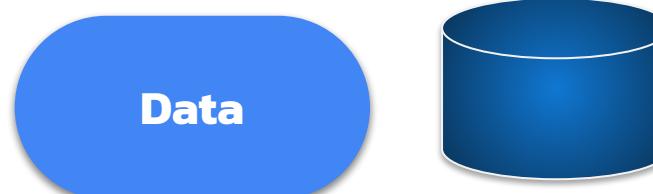
The Art of
Data Modelling

06

What's next?
Quiz, Homework, Next week

DataStax

1. Analyze raw data
2. Identify entities, their properties and relations
3. Design tables, using **normalization** and foreign keys.
4. Use JOIN when doing queries to join normalized data from multiple tables



A diagram illustrating a relational database structure. It shows two tables: 'Employees' and 'Departments'. A line connects the 'departmentId' column in the 'Employees' table to the 'departmentId' column in the 'Departments' table, indicating a foreign key relationship.

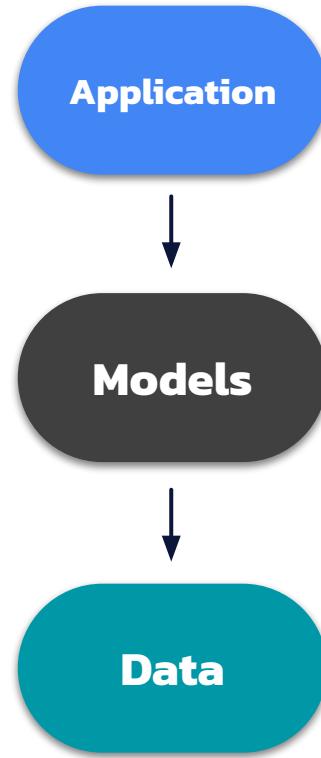
Employees		
userId	firstName	lastName
1	Edgar	Codd
2	Raymond	Boyce

Departments	
departmentId	department
1	Engineering
2	Math

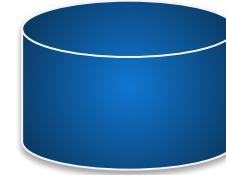


Relational Data Modelling

1. Analyze user behaviour
(customer first!)
2. Identify workflows, their dependencies
and needs
3. Define Queries to fulfill these workflows
4. Knowing the queries, design tables,
using **denormalization**.
5. Use BATCH when inserting or updating
denormalized data of multiple tables

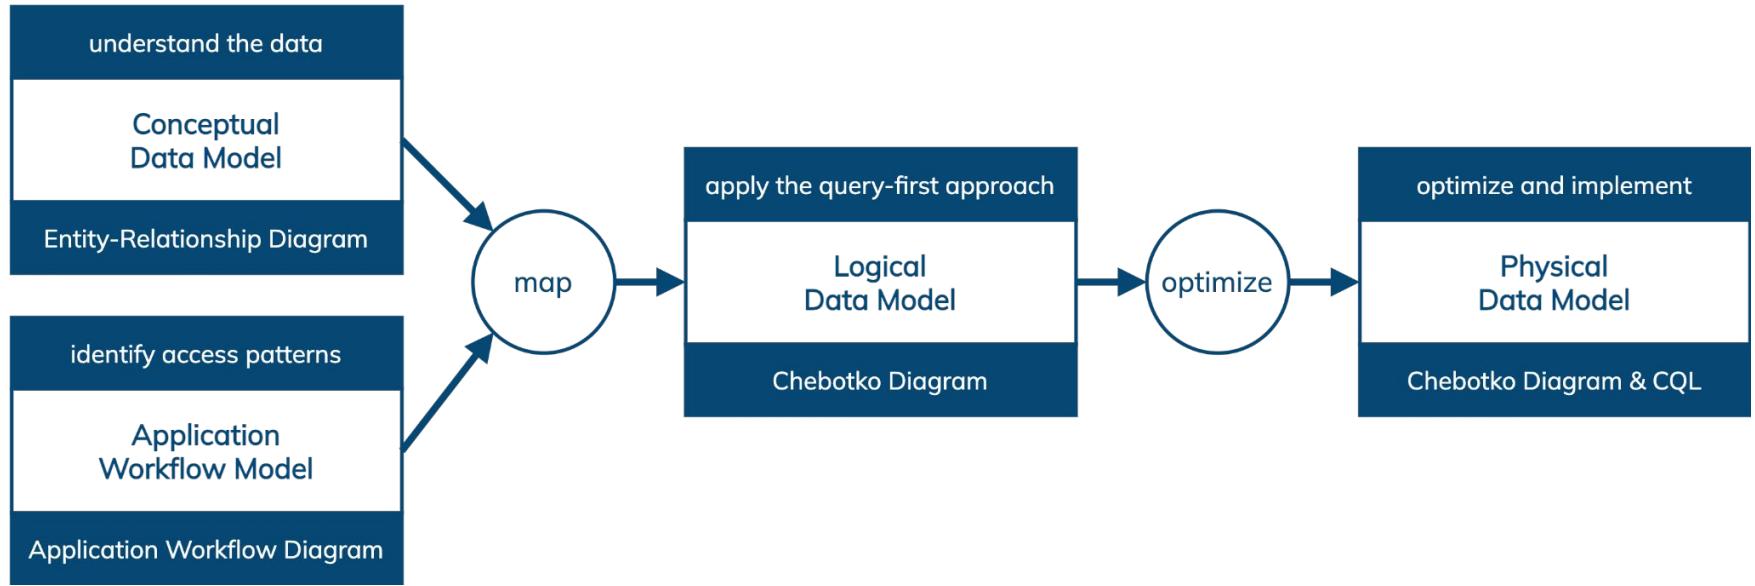


Employees			
userId	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math
3	Sage	Lahja	Math
4	Juniper	Jones	Botany



NoSQL Data Modelling

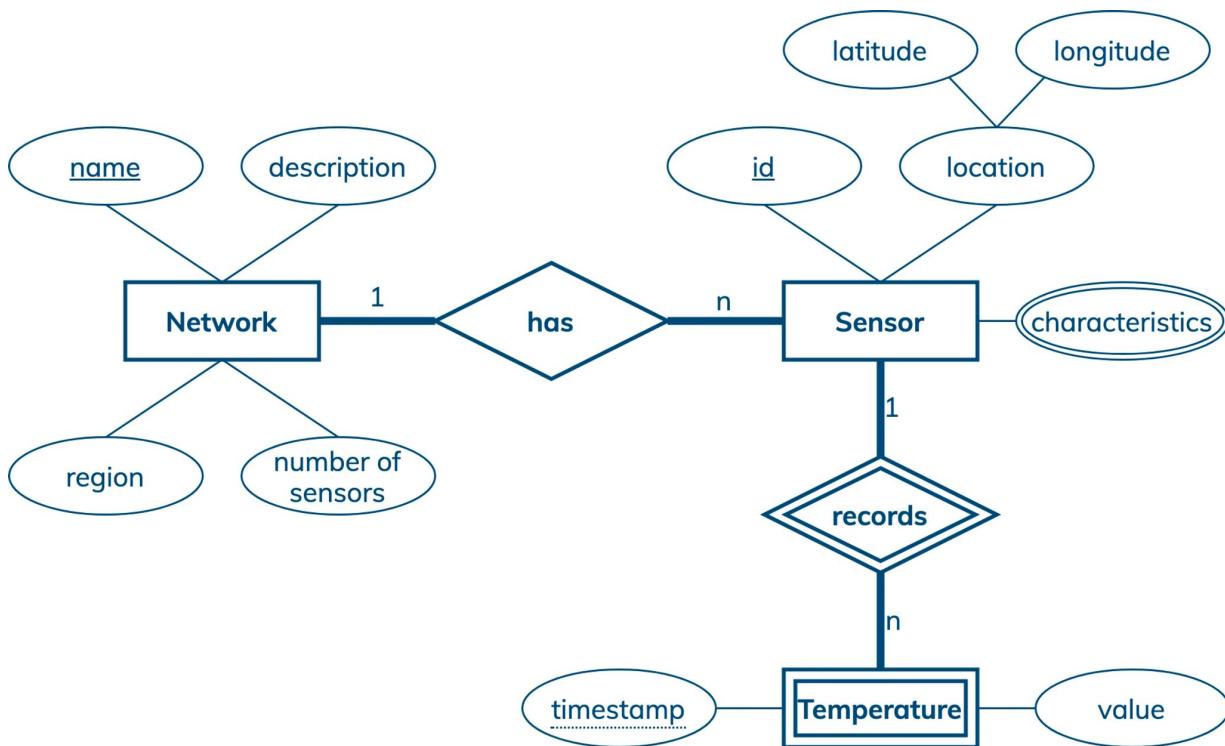




very often, 1 query = 1 table

Data Modelling Methodology





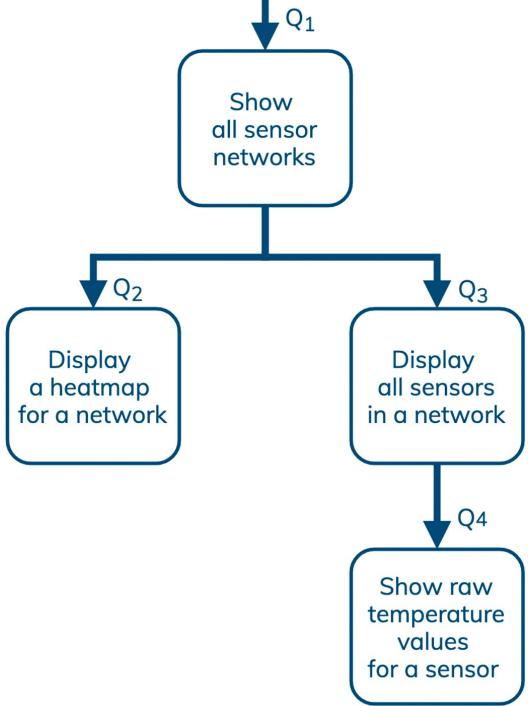
Entity-Relationship Diagram

DataStax

DataStax Developers

Conceptual Data Model

sky



Data access patterns

Q1: Find information about all networks; order by name (asc)

Q2: Find hourly average temperatures for every sensor in a specified network for a given date range; order by date (desc) and hour (desc)

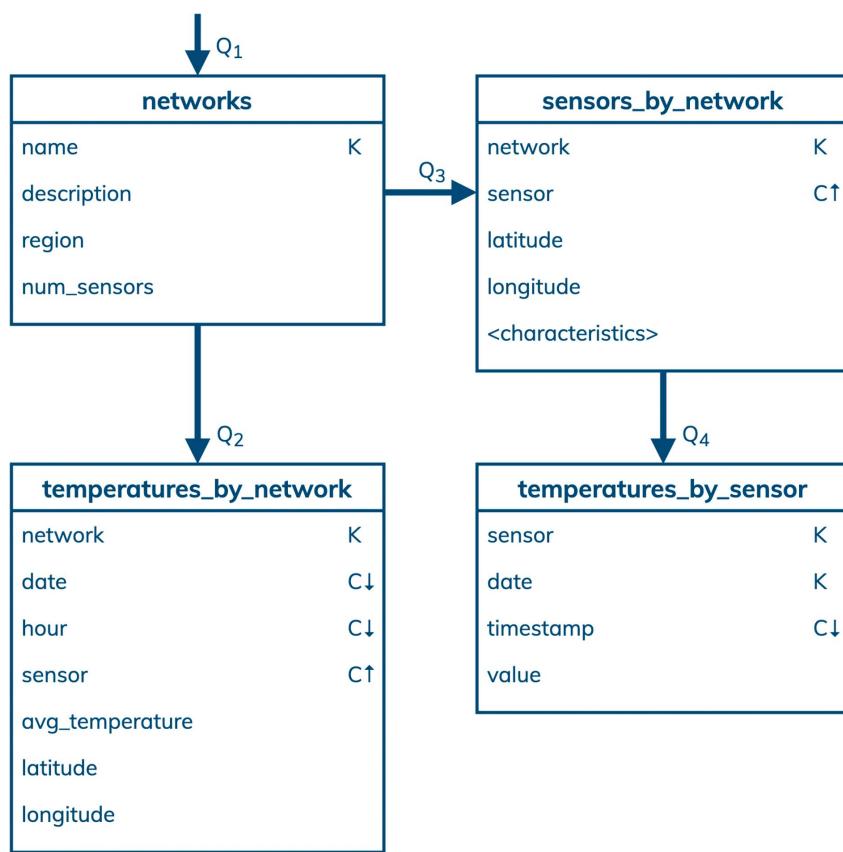
Q3: Find information about all sensors in a specified network

Q4: Find raw measurements for a particular sensor on a specified date; order by timestamp (desc)



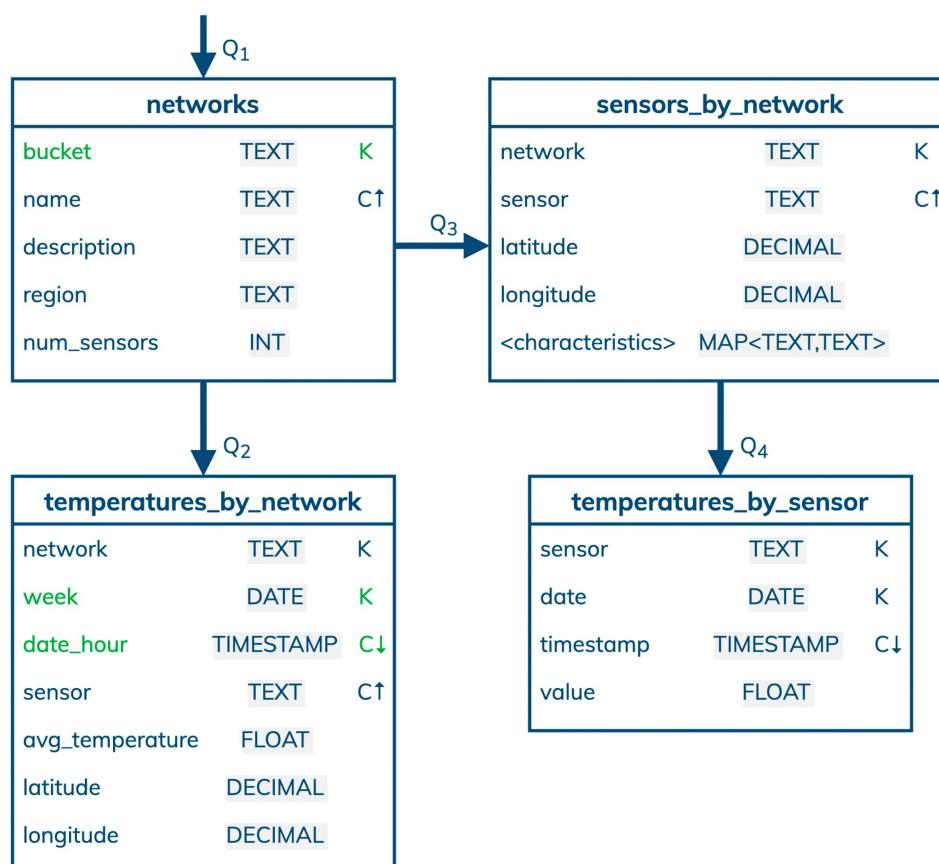
Application Workflow





Chebotko Diagram

Mapping (1 / 2)



Chebotko Diagram

Optimize

```
CREATE TABLE networks (
    bucket TEXT,
    name TEXT,
    description TEXT,
    region TEXT,
    num_sensors INT,
    PRIMARY KEY ((bucket),name)
);
```

↓ Q2

```
CREATE TABLE temperatures_by_network (
    network TEXT,
    week DATE,
    date_hour TIMESTAMP,
    sensor TEXT,
    avg_temperature FLOAT,
    latitude DECIMAL,
    longitude DECIMAL,
    PRIMARY KEY ((network,week),date_hour,sensor)
) WITH CLUSTERING ORDER BY (date_hour DESC, sensor ASC);
```

Q3

```
CREATE TABLE sensors_by_network (
    network TEXT,
    sensor TEXT,
    latitude DECIMAL,
    longitude DECIMAL,
    characteristics MAP<TEXT,TEXT>,
    PRIMARY KEY ((network),sensor)
);
```

Q4

```
CREATE TABLE temperatures_by_sensor (
    sensor TEXT,
    date DATE,
    timestamp TIMESTAMP,
    value FLOAT,
    PRIMARY KEY ((sensor,date),timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);
```

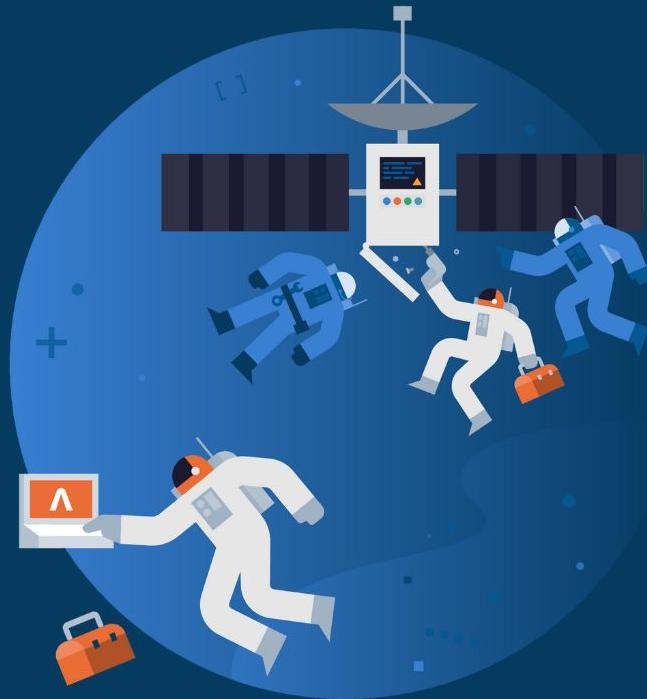


DataStax



sky

Cassandra Query Language



Lab 3

CRUD Operations

[http://github.com/datastaxdevs/
workshop-cassandra-fundamentals](http://github.com/datastaxdevs/workshop-cassandra-fundamentals)



01



Why NoSQL ?

02

Introduction to
Apache Cassandra™

03

Distributed Architecture
The CAP Theorem

04

Tables
and Partitions

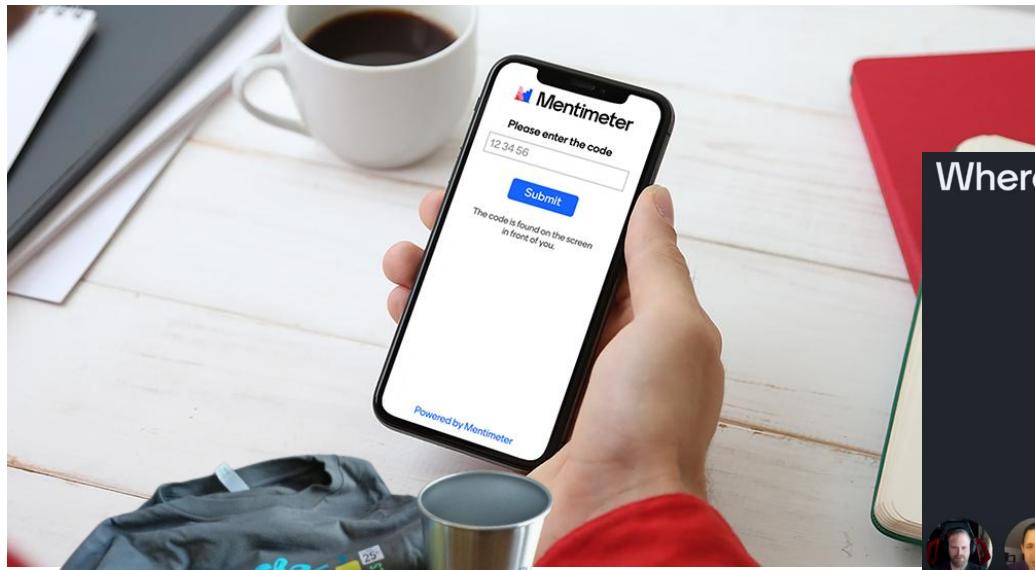
05

The Art of
Data Modelling

06

What's next?
Quiz, Homework, Next week

DataStax



Open a new TAB

CODE= XXXXX



"Menti" for survey and quiz !

Where are you from?



menti.com ⇒ enter code
Don't answer in YT chat
Look at phone (not at YT)
Keep it open for later



SWAG WINNERS



Congratulations to 1st, 2nd and 3rd place on the Menti quiz!

To claim your prize, please send an email to:

gary.harvey@datastax.com

**** Include a screenshot of your Menti screen**



Swag Winners!



Homework

!homework



Intro to Cassandra Homework

Welcome and thank you! Here you can submit your homework for the datastax developers "Intro to Cassandra" workshop. In case of any questions please contact organisers at <https://dtsx.io/aleks> or just send an email to aleksandr.volochnev@datastax.com

- Workshop materials: <https://github.com/datastaxdevs/workshop-intro-to-cassandra>
- Discord chat: <https://dtsx.io/discord>

Switch account 

The name and photo associated with your Google account will be recorded when you upload files and submit this form. Only the email you enter is part of your response.

* Required



DataStax Developers



Join our 17k Discord Community

DataStax Developers



!discord

dtsx.io/discord

A screenshot of the DataStax Developers Discord server. The left sidebar shows various channels: Événements, #moderator-only, WELCOME, start-here, code-of-conduct, introductions, upcoming-events, useful-resources, memes, your-ideas, the-stage, WORKSHOPS, #workshop-chat, #workshop-feedback, workshop-materials, upcoming-workshops, ASTRADB, getting-started, astra-apis, astra-development, sample-applications, and APACHE CASSANDRA. The main channel, #workshop-chat, contains messages from users like RIGGITYREKT, Erick Ramirez, Cedrick Lunven, and others. A video player at the top shows a presentation slide with a network diagram. The right sidebar lists PRESENTER — 1 (David Jones-Giardi), HELPER — 7 (012345, AaronP, Binary, Chelsea Navo, Jeremy Hanna, John Sanda, Patrick_McFadin), and EN LIGNE — 560 (various user names).

Discord Community



Thank You!

